

```
In [82]: %pip install pandas
          %pip install matplotlib
          %pip install seaborn
          %pip install scikit-learn
          %pip install nltk
          %pip install wordcloud
          %pip install emoji
```

Requirement already satisfied: pandas in ./venv/lib/python3.10/site-packages (2.2.3)

Requirement already satisfied: numpy>=1.22.4 in ./venv/lib/python3.10/site-packages (from pandas) (2.2.5)

Requirement already satisfied: python-dateutil>=2.8.2 in ./venv/lib/python3.10/site-packages (from pandas) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.10/site-packages (from pandas) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.10/site-packages (from pandas) (2025.2)

Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)

Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: matplotlib in ./venv/lib/python3.10/site-packages (3.10.1)

Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.10/site-packages (from matplotlib) (1.3.2)

Requirement already satisfied: cycycler>=0.10 in ./venv/lib/python3.10/site-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.10/site-packages (from matplotlib) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib) (1.4.8)

Requirement already satisfied: numpy>=1.23 in ./venv/lib/python3.10/site-packages (from matplotlib) (2.2.5)

Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.10/site-packages (from matplotlib) (25.0)

Requirement already satisfied: pillow>=8 in ./venv/lib/python3.10/site-packages (from matplotlib) (11.2.1)

Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in ./venv/lib/python3.10/site-packages (from matplotlib) (2.9.0.post0)

Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib) (1.17.0)

Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: seaborn in ./venv/lib/python3.10/site-packages (0.13.2)

Requirement already satisfied: numpy!=1.24.0,>=1.20 in ./venv/lib/python3.10/site-packages (from seaborn) (2.2.5)

Requirement already satisfied: pandas>=1.2 in ./venv/lib/python3.10/site-packages (from seaborn) (2.2.3)

Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in ./venv/lib/python3.10/site-packages (from seaborn) (3.10.1)

Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.2)

Requirement already satisfied: cycycler>=0.10 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.8)

Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (25.0)

Requirement already satisfied: pillow>=8 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.2.1)

Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in ./venv/lib/python3.10/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.10/site-packages (from pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.10/site-packages (from pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.17.0)

Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: scikit-learn in ./venv/lib/python3.10/site-packages (1.6.1)

Requirement already satisfied: numpy>=1.19.5 in ./venv/lib/python3.10/site-packages (from scikit-learn) (2.2.5)

Requirement already satisfied: scipy>=1.6.0 in ./venv/lib/python3.10/site-packages (from scikit-learn) (1.15.2)

Requirement already satisfied: joblib>=1.2.0 in ./venv/lib/python3.10/site-packages (from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in ./venv/lib/python3.10/site-packages (from scikit-learn) (3.6.0)

Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: nltk in ./venv/lib/python3.10/site-packages (3.9.1)

Requirement already satisfied: click in ./venv/lib/python3.10/site-packages (from nltk) (8.1.8)

Requirement already satisfied: joblib in ./venv/lib/python3.10/site-packages (from nltk) (1.4.2)

Requirement already satisfied: regex>=2021.8.3 in ./venv/lib/python3.10/site-packages (from nltk) (2024.11.6)

Requirement already satisfied: tqdm in ./venv/lib/python3.10/site-packages (from nltk) (4.67.1)

Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: wordcloud in ./venv/lib/python3.10/site-packages (1.9.4)

Requirement already satisfied: numpy>=1.6.1 in ./venv/lib/python3.10/site-packages (from wordcloud) (2.2.5)

Requirement already satisfied: pillow in ./venv/lib/python3.10/site-packages (from wordcloud) (11.2.1)

Requirement already satisfied: matplotlib in ./venv/lib/python3.10/site-packages (from wordcloud) (3.10.1)

Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (1.3.2)

Requirement already satisfied: cycler>=0.10 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (1.4.8)

Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (25.0)

Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in ./venv/lib/python3.10/site-packages (from matplotlib->wordcloud) (2.9.0.post0)

Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib->wordcloud) (1.17.0)

Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: emoji in ./venv/lib/python3.10/site-packages (2.14.1)

Note: you may need to restart the kernel to use updated packages.

```
In [83]: import matplotlib.pyplot as plt
```

```
import seaborn as sns
import numpy as np

import pandas as pd
df = pd.read_csv("scitweets_export.tsv", sep="\t")
df.head()
```

Out[83]:

	Unnamed: 0	tweet_id	text	science_related	scientific_claim	scie
0	0	316669998137483264	Knees are a bit sore. i guess that's a sign th...	0	0.0	
1	1	319090866545385472	McDonald's breakfast stop then the gym 🏀💪	0	0.0	
2	2	322030931022065664	Can any Gynecologist with Cancer Experience ex...	1	1.0	
3	3	322694830620807168	Couch-lock highs lead to sleeping in the couch...	1	1.0	
4	4	328524426658328576	Does daily routine help prevent problems with ...	1	1.0	

In [120... *# Remove warning messages*

```
import warnings

from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore", category=UserWarning, module="matplotlib")
warnings.filterwarnings("ignore", category=UserWarning, module="seaborn")
warnings.filterwarnings("ignore", category=UserWarning, module="pandas")
warnings.filterwarnings("ignore", category=UserWarning, module="wordcloud")
warnings.filterwarnings("ignore", category=FutureWarning, module="pandas")
warnings.filterwarnings("ignore", category=DeprecationWarning, module="pa")
warnings.filterwarnings("ignore", category=ConvergenceWarning, module="sk
```

In [85]:

```
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(stop_words='english', max_features=5000)
X_all = vectorizer.fit_transform(df['text'])
```

In [119... *# Compare different feature extraction methods*

```
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_selection import SelectKBest, chi2, mutual_info_clas

# Create label for Task 1
y_task1 = df['science_related']
```

```
# 1. Count Vectorizer
count_vec = CountVectorizer(max_features=5000)
X_count = count_vec.fit_transform(df['text'])

# 2. TF-IDF with more parameters
tfidf_vec = TfidfVectorizer(max_features=5000,
                             min_df=5,
                             max_df=0.8,
                             ngram_range=(1, 2))
X_tfidf = tfidf_vec.fit_transform(df['text'])

# 3. TF-IDF with preprocessing already done
tfidf_processed = TfidfVectorizer(max_features=5000)
X_tfidf_processed = tfidf_processed.fit_transform(df['text'])

# Compare feature extraction methods
print(f"Count Vectorizer Features: {X_count.shape}")
print(f"TF-IDF Vectorizer Features: {X_tfidf.shape}")
print(f"TF-IDF on Preprocessed Text Features: {X_tfidf_processed.shape}")

# Feature selection using Chi-squared
selector_chi2 = SelectKBest(chi2, k=100)
X_chi2 = selector_chi2.fit_transform(X_tfidf, y_task1)

# Feature selection using Mutual Information
selector_mi = SelectKBest(mutual_info_classif, k=100)
X_mi = selector_mi.fit_transform(X_tfidf, y_task1)

print(f"\nFeatures after Chi-squared selection: {X_chi2.shape}")
print(f"Features after Mutual Information selection: {X_mi.shape}")

# Get and visualize the most important features
chi2_selected_indices = selector_chi2.get_support(indices=True)
mi_selected_indices = selector_mi.get_support(indices=True)

chi2_feature_names = np.array(tfidf_vec.get_feature_names_out())[chi2_selected_indices]
mi_feature_names = np.array(tfidf_vec.get_feature_names_out())[mi_selected_indices]

# Plot top 20 features by importance
plt.figure(figsize=(16, 8))

plt.subplot(1, 2, 1)
chi2_scores = selector_chi2.scores_[chi2_selected_indices]
chi2_features_df = pd.DataFrame({'Feature': chi2_feature_names, 'Score': chi2_scores})
chi2_features_df = chi2_features_df.sort_values('Score', ascending=False)
sns.barplot(x='Score', y='Feature', data=chi2_features_df)
plt.title('Top 20 Features - Chi-squared')

plt.subplot(1, 2, 2)
mi_scores = selector_mi.scores_[mi_selected_indices]
mi_features_df = pd.DataFrame({'Feature': mi_feature_names, 'Score': mi_scores})
mi_features_df = mi_features_df.sort_values('Score', ascending=False).head(20)
sns.barplot(x='Score', y='Feature', data=mi_features_df)
plt.title('Top 20 Features - Mutual Information')

plt.tight_layout()
plt.show()

# We'll use the TF-IDF on preprocessed text for subsequent modeling
X_selected = X_tfidf
```


Count Vectorizer Features: (1140, 5000)

TF-IDF Vectorizer Features: (1140, 693)

TF-IDF on Preprocessed Text Features: (1140, 5000)

Features after Chi-squared selection: (1140, 100)

Features after Mutual Information selection: (1140, 100)



```
In [87]: # Create labels for all tasks
df['task1_label'] = df['science_related']

df_sci = df[df['science_related'] == 1].copy()
df_sci['task2_label'] = ((df_sci['scientific_claim'] == 1.0) | (df_sci['s
df_sci['task3_label'] = df_sci[['scientific_claim', 'scientific_reference
df_sci['task3_label'] = df_sci['task3_label'].map({
    'scientific_claim': 0,
    'scientific_reference': 1,
    'scientific_context': 2
})
```

```
In [88]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB, MultinomialNB, ComplementNB,
from sklearn.metrics import classification_report, confusion_matrix, accu
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

# Function to evaluate model performance
def evaluate_model(model, X_train, X_test, y_train, y_test, model_name):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)
    report = classification_report(y_test, y_pred, output_dict=True)
    cm = confusion_matrix(y_test, y_pred)

    print(f"Model: {model_name}")
    print(f"Accuracy: {accuracy:.4f}")
    print(classification_report(y_test, y_pred))

    return {
        'model_name': model_name,
```

```
        'accuracy': accuracy,
        'report': report,
        'confusion_matrix': cm,
        'y_pred': y_pred
    }

    # Get a dense version of our features for Gaussian NB
    X_dense = X_selected.toarray()

    # Split data for Task 1
    X_train_task1, X_test_task1, y_train_task1, y_test_task1 = train_test_split(
        X_dense, df['task1_label'], test_size=0.2, random_state=42
    )

    # Compare different NB variants for Task 1
    nb_models = {
        'Gaussian NB': GaussianNB(),
        'Multinomial NB': MultinomialNB(),
        'Complement NB': ComplementNB(),
        'Bernoulli NB': BernoulliNB(),
        'KNN' : KNeighborsClassifier(n_neighbors=5),
        'SVM' : SVC(kernel='linear', C=1),
        'Logistic Regression': LogisticRegression(max_iter=1000),
        'Random Forest': RandomForestClassifier(n_estimators=100),
    }

    task1_results = {}
    print("Task 1: Science Related Classification\n" + "="*40)
    for name, model in nb_models.items():
        task1_results[name] = evaluate_model(
            model, X_train_task1, X_test_task1, y_train_task1, y_test_task1,
        )
    print("\n")
```

## Task 1: Science Related Classification

=====

Model: Gaussian NB

Accuracy: 0.6711

	precision	recall	f1-score	support
0	0.79	0.66	0.72	146
1	0.53	0.70	0.60	82
accuracy			0.67	228
macro avg	0.66	0.68	0.66	228
weighted avg	0.70	0.67	0.68	228

Model: Multinomial NB

Accuracy: 0.7544

	precision	recall	f1-score	support
0	0.74	0.96	0.83	146
1	0.84	0.39	0.53	82
accuracy			0.75	228
macro avg	0.79	0.67	0.68	228
weighted avg	0.77	0.75	0.73	228

Model: Complement NB

Accuracy: 0.7851

	precision	recall	f1-score	support
0	0.84	0.82	0.83	146
1	0.69	0.73	0.71	82
accuracy			0.79	228
macro avg	0.77	0.77	0.77	228
weighted avg	0.79	0.79	0.79	228

Model: Bernoulli NB

Accuracy: 0.7895

	precision	recall	f1-score	support
0	0.80	0.90	0.85	146
1	0.77	0.59	0.67	82
accuracy			0.79	228
macro avg	0.78	0.74	0.76	228
weighted avg	0.79	0.79	0.78	228

Model: KNN

Accuracy: 0.7105

	precision	recall	f1-score	support
0	0.80	0.73	0.76	146
1	0.58	0.68	0.63	82



accuracy			0.71	228
macro avg	0.69	0.70	0.70	228
weighted avg	0.72	0.71	0.71	228

Model: SVM

Accuracy: 0.7895

	precision	recall	f1-score	support
0	0.79	0.92	0.85	146
1	0.79	0.56	0.66	82

accuracy			0.79	228
macro avg	0.79	0.74	0.75	228
weighted avg	0.79	0.79	0.78	228

Model: Logistic Regression

Accuracy: 0.7675

	precision	recall	f1-score	support
0	0.74	0.97	0.84	146
1	0.89	0.40	0.55	82

accuracy			0.77	228
macro avg	0.82	0.69	0.70	228
weighted avg	0.80	0.77	0.74	228

Model: Random Forest

Accuracy: 0.7544

	precision	recall	f1-score	support
0	0.74	0.94	0.83	146
1	0.80	0.43	0.56	82

accuracy			0.75	228
macro avg	0.77	0.68	0.69	228
weighted avg	0.76	0.75	0.73	228

```
In [89]: from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score

# Initialize k-fold cross validation
k_folds = 10
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)

# Dictionary to store results
cv_results = {}

# Perform k-fold cross validation for each model
for model_name, model in nb_models.items():
    print(f"Performing {k_folds}-fold cross-validation for {model_name}..")
```

```
# Initialize lists to store performance metrics for each fold
fold_accuracy = []
fold_precision = []
fold_recall = []
fold_f1 = []

# For each fold
for fold, (train_idx, test_idx) in enumerate(skf.split(X_dense, y_task1)):
    # Split data
    X_train_fold, X_test_fold = X_dense[train_idx], X_dense[test_idx]
    y_train_fold, y_test_fold = y_task1.iloc[train_idx], y_task1.iloc[test_idx]

    # Train model
    model.fit(X_train_fold, y_train_fold)

    # Make predictions
    y_pred_fold = model.predict(X_test_fold)

    # Calculate metrics
    acc = accuracy_score(y_test_fold, y_pred_fold)
    prec = precision_score(y_test_fold, y_pred_fold, zero_division=0)
    rec = recall_score(y_test_fold, y_pred_fold, zero_division=0)
    f1 = f1_score(y_test_fold, y_pred_fold, zero_division=0)

    fold_accuracy.append(acc)
    fold_precision.append(prec)
    fold_recall.append(rec)
    fold_f1.append(f1)

# Store average metrics and standard deviations
cv_results[model_name] = {
    'accuracy': {
        'mean': np.mean(fold_accuracy),
        'std': np.std(fold_accuracy)
    },
    'precision': {
        'mean': np.mean(fold_precision),
        'std': np.std(fold_precision)
    },
    'recall': {
        'mean': np.mean(fold_recall),
        'std': np.std(fold_recall)
    },
    'f1': {
        'mean': np.mean(fold_f1),
        'std': np.std(fold_f1)
    }
}

print(f" Average: Accuracy={cv_results[model_name]['accuracy']['mean']}")
print(f" Average: F1 Score={cv_results[model_name]['f1']['mean']:.4f}")
print()

# Create DataFrame for visualization
results_df = pd.DataFrame({
    'Model': [],
    'Metric': [],
    'Mean': [],
    'Std': []
})
```

```

}))

for model_name in cv_results:
    for metric in ['accuracy', 'precision', 'recall', 'f1']:
        results_df = pd.concat([results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': [metric.capitalize()],
            'Mean': [cv_results[model_name][metric]['mean']],
            'Std': [cv_results[model_name][metric]['std']]
        })], ignore_index=True)

# Sort models by accuracy
model_order = results_df[results_df['Metric'] == 'Accuracy'].sort_values(

# Create plots
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")

# Create bar plot for accuracy
ax = sns.barplot(
    data=results_df[results_df['Metric'] == 'Accuracy'],
    x='Model',
    y='Mean',
    order=model_order,
    palette='Blues_d'
)

# Add error bars
for i, model in enumerate(model_order):
    row = results_df[(results_df['Model'] == model) & (results_df['Metric'] == 'Accuracy')]
    ax.errorbar(
        i, row['Mean'], yerr=row['Std'],
        fmt='o', color='black', elinewidth=2, capsize=6
    )

# Add value labels on top of bars
for i, bar in enumerate(ax.patches):
    ax.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height() + 0.01,
        f"{bar.get_height():.3f}",
        ha='center',
        fontsize=10
    )

plt.title(f'Model Accuracy Comparison with {k_folds}-Fold Cross Validation')
plt.xlabel('Model', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.ylim([0.5, 1.0])
plt.tight_layout()
plt.show()

# Create a grouped bar chart for all metrics
plt.figure(figsize=(15, 10))
sns.set_style("whitegrid")

# Create grouped bar plot
ax = sns.barplot(
    data=results_df,

```

```

        x='Model',
        y='Mean',
        hue='Metric',
        order=model_order,
        palette='Set2'
    )

plt.title(f'Model Performance Metrics Comparison ({k_folds}-Fold CV)', fo
plt.xlabel('Model', fontsize=14)
plt.ylabel('Score', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Metric', loc='lower right')
plt.ylim([0.5, 1.0])
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```

Performing 10-fold cross-validation for Gaussian NB...

Average: Accuracy=0.6728 ± 0.0245

Average: F1 Score=0.6090 ± 0.0290

Performing 10-fold cross-validation for Multinomial NB...

Average: Accuracy=0.7868 ± 0.0232

Average: F1 Score=0.5587 ± 0.0653

Performing 10-fold cross-validation for Complement NB...

Average: Accuracy=0.7772 ± 0.0229

Average: F1 Score=0.6921 ± 0.0204

Performing 10-fold cross-validation for Bernoulli NB...

Average: Accuracy=0.7982 ± 0.0215

Average: F1 Score=0.6795 ± 0.0355

Performing 10-fold cross-validation for KNN...

Average: Accuracy=0.7114 ± 0.0367

Average: F1 Score=0.5443 ± 0.0647

Performing 10-fold cross-validation for SVM...

Average: Accuracy=0.7930 ± 0.0319

Average: F1 Score=0.6562 ± 0.0459

Performing 10-fold cross-validation for Logistic Regression...

Average: Accuracy=0.7904 ± 0.0262

Average: F1 Score=0.5900 ± 0.0711

Performing 10-fold cross-validation for Random Forest...



Average: Accuracy=0.7807 ± 0.0222

Average: F1 Score=0.5803 ± 0.0515

/tmp/ipykernel\_368067/4283035198.py:93: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(
```



## Parameter optimization

```
In [90]: from sklearn.model_selection import GridSearchCV
```

```
for model_name, model in nb_models.items():
```

```
print(f"Performing Grid Search for {model_name}...")
# Define the parameter grid
if model_name == 'Gaussian NB':
    param_grid = {
        'var_smoothing': np.
        logspace(0, -9, num=100)
    }
elif model_name == 'Multinomial NB':
    param_grid = {
        'alpha': np.logspace(-3, 3, num=100),
        'fit_prior': [True, False]
    }
elif model_name == 'Complement NB':
    param_grid = {
        'alpha': np.logspace(-3, 3, num=100),
        'fit_prior': [True, False]
    }
elif model_name == 'Bernoulli NB':
    param_grid = {
        'alpha': np.logspace(-3, 3, num=100),
        'fit_prior': [True, False]
    }
elif model_name == 'KNN':
    param_grid = {
        'n_neighbors': [3, 5, 7, 9],
        'weights': ['uniform', 'distance'],
        'metric': ['euclidean', 'manhattan']
    }
elif model_name == 'SVM':
    param_grid = {
        'C': np.logspace(-3, 3, num=100),
        'kernel': ['linear', 'rbf'],
        'gamma': ['scale', 'auto']
    }
elif model_name == 'Logistic Regression':
    param_grid = {
        'C': np.logspace(-3, 3, num=100),
        'penalty': ['l1', 'l2'],
        'solver': ['liblinear', 'saga']
    }
elif model_name == 'Random Forest':
    param_grid = {
        'n_estimators': [50, 100, 200],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    }
else:
    continue

grid_search = GridSearchCV(
    estimator=model,
    param_grid=param_grid,
    scoring='accuracy',
    cv=StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=4
    n_jobs=-1,
    verbose=1
)
```

```
# Fit the grid search to the data
grid_search.fit(X_train_task1, y_train_task1)

# Get the best parameters and best score
print(f"Best Parameters: {grid_search.best_params_}")
print(f"Best Cross-Validation Score: {grid_search.best_score_:.4f}")

# Evaluate the best model
best_model = grid_search.best_estimator_
best_model_accuracy = best_model.score(X_test_task1, y_test_task1)
print(f"Test Accuracy with Best Parameters: {best_model_accuracy:.4f}")

# Visualize parameter impact
results = pd.DataFrame(grid_search.cv_results_)

# Plot effect of parameters

if model_name == 'Gaussian NB':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_var_smoothing', y='mean_test_
    plt.xscale('log')
    plt.title(f'Gaussian NB: Effect of var_smoothing on Accuracy')
    plt.xlabel('var_smoothing (log scale)')
    plt.ylabel('Mean Test Score')
    plt.grid(True)
    plt.show()
elif model_name == 'Multinomial NB':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_alpha', y='mean_test_score',
    plt.xscale('log')
    plt.title(f'Multinomial NB: Effect of alpha and fit_prior on Accu
    plt.xlabel('alpha (log scale)')
    plt.ylabel('Mean Test Score')
    plt.grid(True)
    plt.show()
elif model_name == 'Complement NB':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_alpha', y='mean_test_score',
    plt.xscale('log')
    plt.title(f'Complement NB: Effect of alpha and fit_prior on Accur
    plt.xlabel('alpha (log scale)')
    plt.ylabel('Mean Test Score')
    plt.grid(True)
    plt.show()
elif model_name == 'Bernoulli NB':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_alpha', y='mean_test_score',
    plt.xscale('log')
    plt.title(f'Bernoulli NB: Effect of alpha and fit_prior on Accura
    plt.xlabel('alpha (log scale)')
    plt.ylabel('Mean Test Score')
    plt.grid(True)
    plt.show()
elif model_name == 'KNN':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_n_neighbors', y='mean_test_sc
    plt.title(f'KNN: Effect of n_neighbors and weights on Accuracy')
    plt.xlabel('n_neighbors')
    plt.ylabel('Mean Test Score')
    plt.grid(True)
```

```

plt.show()
elif model_name == 'SVM':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_C', y='mean_test_score', hue=
plt.xscale('log')
plt.title(f'SVM: Effect of C and kernel on Accuracy')
plt.xlabel('C (log scale)')
plt.ylabel('Mean Test Score')
plt.grid(True)
plt.show()
elif model_name == 'Logistic Regression':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_C', y='mean_test_score', hue=
plt.xscale('log')
plt.title(f'Logistic Regression: Effect of C and penalty on Accur
plt.xlabel('C (log scale)')
plt.ylabel('Mean Test Score')
plt.grid(True)
plt.show()
elif model_name == 'Random Forest':
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=results, x='param_n_estimators', y='mean_test_s
plt.title(f'Random Forest: Effect of n_estimators and max_depth o
plt.xlabel('n_estimators')
plt.ylabel('Mean Test Score')
plt.grid(True)
plt.show()
else:
    print(f"No visualization available for {model_name} model.")

```


Performing Grid Search for Gaussian NB...

Fitting 10 folds for each of 100 candidates, totalling 1000 fits

Best Parameters: {'var\_smoothing': np.float64(0.01519911082952933)}

Best Cross-Validation Score: 0.7500

Test Accuracy with Best Parameters: 0.7325






Performing Grid Search for Multinomial NB...

Fitting 10 folds for each of 200 candidates, totalling 2000 fits

Best Parameters: {'alpha': np.float64(0.4641588833612782), 'fit\_prior': True}

Best Cross-Validation Score: 0.8093

Test Accuracy with Best Parameters: 0.7982




Performing Grid Search for Complement NB...

Fitting 10 folds for each of 200 candidates, totalling 2000 fits

Best Parameters: {'alpha': np.float64(2.848035868435802), 'fit\_prior': True}

Best Cross-Validation Score: 0.8060

Test Accuracy with Best Parameters: 0.7895




Performing Grid Search for Bernoulli NB...

Fitting 10 folds for each of 200 candidates, totalling 2000 fits

Best Parameters: {'alpha': np.float64(0.1747528400007685), 'fit\_prior': True}

Best Cross-Validation Score: 0.8224

Test Accuracy with Best Parameters: 0.8070




Performing Grid Search for KNN...

Fitting 10 folds for each of 16 candidates, totalling 160 fits

Best Parameters: {'metric': 'euclidean', 'n\_neighbors': 7, 'weights': 'uniform'}

Best Cross-Validation Score: 0.7412

Test Accuracy with Best Parameters: 0.7237




Performing Grid Search for SVM...

Fitting 10 folds for each of 400 candidates, totalling 4000 fits

Best Parameters: {'C': np.float64(1.4174741629268048), 'gamma': 'scale', 'kernel': 'rbf'}

Best Cross-Validation Score: 0.8049

Test Accuracy with Best Parameters: 0.7719



Performing Grid Search for Logistic Regression...

Fitting 10 folds for each of 400 candidates, totalling 4000 fits

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]



[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]



[illegible]

[illegible]

[illegible]

[illegible]

[illegible]


[illegible]

[illegible]

```

/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sk
learn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reach
ed which means the coef_ did not converge
  warnings.warn(
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sk
learn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reach
ed which means the coef_ did not converge
  warnings.warn(
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sk
learn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reach
ed which means the coef_ did not converge
  warnings.warn(
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sk
learn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reach
ed which means the coef_ did not converge
  warnings.warn(
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sk
learn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reach
ed which means the coef_ did not converge
  warnings.warn(
Best Parameters: {'C': np.float64(4.328761281083062), 'penalty': 'l1', 'sol
ver': 'saga'}
Best Cross-Validation Score: 0.8081
Test Accuracy with Best Parameters: 0.7939

```




```

Performing Grid Search for Random Forest...
Fitting 10 folds for each of 108 candidates, totalling 1080 fits
Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_s
plit': 10, 'n_estimators': 50}
Best Cross-Validation Score: 0.7818
Test Accuracy with Best Parameters: 0.7675

```





## Training comparison with best parameter

```
In [91]: from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, precision_score, recall_score

# Create dictionary of models with their best parameters based on grid se

# Define models with their best parameters based on grid search results
best_models = {
    'Gaussian NB': GaussianNB(var_smoothing=0.0152), #ample value - updat
    'Multinomial NB': MultinomialNB(alpha=0.4641, fit_prior=True), # Exa
    'Complement NB': ComplementNB(alpha=2.8480, fit_prior=True), # Examp
    'Bernoulli NB': BernoulliNB(alpha=0.17475, fit_prior=True), # Examp
    'KNN': KNeighborsClassifier(n_neighbors=7, weights='uniform', metric=
    'SVM': SVC(C=1.4174742, kernel='rbf', gamma='scale', probability=True
    'Logistic Regression': LogisticRegression(C=np.float64(4.328761281083
    'Random Forest': RandomForestClassifier(n_estimators=200, max_depth=N
}

# Dictionary to store results
best_cv_results = {}

# Perform k-fold cross validation for each model
for model_name, model in best_models.items():
    print(f"Performing {k_folds}-fold cross-validation for {model_name} w

    # Initialize lists to store performance metrics for each fold
    fold_accuracy = []
    fold_precision = []
    fold_recall = []
    fold_f1 = []
    fold_auc = []
```

```

# For each fold
for fold, (train_idx, test_idx) in enumerate(skf.split(X_dense, y_tas
# Split data
X_train_fold, X_test_fold = X_dense[train_idx], X_dense[test_idx]
y_train_fold, y_test_fold = y_task1.iloc[train_idx], y_task1.iloc

# Train model
model.fit(X_train_fold, y_train_fold)

# Make predictions
y_pred_fold = model.predict(X_test_fold)

# Calculate metrics
acc = accuracy_score(y_test_fold, y_pred_fold)
prec = precision_score(y_test_fold, y_pred_fold, zero_division=0)
rec = recall_score(y_test_fold, y_pred_fold, zero_division=0)
f1 = f1_score(y_test_fold, y_pred_fold, zero_division=0)

# For AUC, we need probability estimates
try:
    if hasattr(model, "predict_proba"):
        y_prob = model.predict_proba(X_test_fold)[: , 1]
        auc_score = roc_auc_score(y_test_fold, y_prob)
        fold_auc.append(auc_score)
except:
    pass

fold_accuracy.append(acc)
fold_precision.append(prec)
fold_recall.append(rec)
fold_f1.append(f1)

# Store average metrics and standard deviations
best_cv_results[model_name] = {
    'accuracy': {
        'mean': np.mean(fold_accuracy),
        'std': np.std(fold_accuracy)
    },
    'precision': {
        'mean': np.mean(fold_precision),
        'std': np.std(fold_precision)
    },
    'recall': {
        'mean': np.mean(fold_recall),
        'std': np.std(fold_recall)
    },
    'f1': {
        'mean': np.mean(fold_f1),
        'std': np.std(fold_f1)
    }
}

if fold_auc:
    best_cv_results[model_name]['auc'] = {
        'mean': np.mean(fold_auc),
        'std': np.std(fold_auc)
    }
    print(f" Average: AUC={best_cv_results[model_name]['auc']['mean']

print(f" Average: Accuracy={best_cv_results[model_name]['accuracy']['

```

```

    print(f"   Average: F1 Score={best_cv_results[model_name]['f1']['mean']}")
    print()

# Create DataFrame for visualization
best_results_df = pd.DataFrame({
    'Model': [],
    'Metric': [],
    'Mean': [],
    'Std': []
})

for model_name in best_cv_results:
    for metric in ['accuracy', 'precision', 'recall', 'f1']:
        best_results_df = pd.concat([best_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': [metric.capitalize()],
            'Mean': [best_cv_results[model_name][metric]['mean']],
            'Std': [best_cv_results[model_name][metric]['std']]
        })], ignore_index=True)
    if 'auc' in best_cv_results[model_name]:
        best_results_df = pd.concat([best_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': ['AUC'],
            'Mean': [best_cv_results[model_name]['auc']['mean']],
            'Std': [best_cv_results[model_name]['auc']['std']]
        })], ignore_index=True)

# Sort models by accuracy
best_model_order = best_results_df[best_results_df['Metric'] == 'Accuracy']

# Create plots
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")

# Create bar plot for accuracy
ax = sns.barplot(
    data=best_results_df[best_results_df['Metric'] == 'Accuracy'],
    x='Model',
    y='Mean',
    order=best_model_order,
    palette='Blues_d'
)

# Add error bars
for i, model in enumerate(best_model_order):
    row = best_results_df[(best_results_df['Model'] == model) & (best_results_df['Metric'] == 'Accuracy')]
    ax.errorbar(
        i, row['Mean'], yerr=row['Std'],
        fmt='o', color='black', elinewidth=2, capsize=6
    )

# Add value labels on top of bars
for i, bar in enumerate(ax.patches):
    ax.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height() + 0.01,
        f"{bar.get_height():.3f}",
        ha='center',
        fontsize=10
    )

```

```

plt.title(f'Model Accuracy Comparison with {k_folds}-Fold Cross Validatio
plt.xlabel('Model', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.ylim([0.5, 1.0])
plt.tight_layout()
plt.show()

# Create a grouped bar chart for all metrics
plt.figure(figsize=(15, 10))
sns.set_style("whitegrid")

# Create grouped bar plot
ax = sns.barplot(
    data=best_results_df[best_results_df['Metric'] != 'AUC'],
    x='Model',
    y='Mean',
    hue='Metric',
    order=best_model_order,
    palette='Set2'
)

plt.title(f'Model Performance Metrics Comparison with Best Parameters ({k
plt.xlabel('Model', fontsize=14)
plt.ylabel('Score', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Metric', loc='lower right')
plt.ylim([0.5, 1.0])
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Create a comparison table between original results and optimized result
comparison_df = pd.DataFrame({
    'Model': [],
    'Metric': [],
    'Original Mean': [],
    'Original Std': [],
    'Optimized Mean': [],
    'Optimized Std': [],
    'Improvement': []
})

for model_name in best_cv_results:
    if model_name in cv_results:
        for metric in ['accuracy', 'precision', 'recall', 'f1']:
            orig_mean = cv_results[model_name][metric]['mean']
            orig_std = cv_results[model_name][metric]['std']
            opt_mean = best_cv_results[model_name][metric]['mean']
            opt_std = best_cv_results[model_name][metric]['std']
            improvement = ((opt_mean - orig_mean) / orig_mean) * 100

            comparison_df = pd.concat([comparison_df, pd.DataFrame({
                'Model': [model_name],
                'Metric': [metric.capitalize()],
                'Original Mean': [orig_mean],
                'Original Std': [orig_std],
                'Optimized Mean': [opt_mean],
                'Optimized Std': [opt_std],

```

```
        'Improvement': [improvement]
    }]], ignore_index=True)

# Plot the improvement in accuracy
plt.figure(figsize=(15, 10))
sns.set_style("whitegrid")

# Filter for accuracy only
acc_comparison = comparison_df[comparison_df['Metric'] == 'Accuracy']
acc_comparison = acc_comparison.sort_values('Improvement', ascending=False)

# Create bar chart of improvements
ax = sns.barplot(
    data=acc_comparison,
    x='Model',
    y='Improvement',
    palette='RdYlGn',
    dodge=False
)

plt.axhline(y=0, color='black', linestyle='--', alpha=0.3)

# Add value labels on bars
for i, bar in enumerate(ax.patches):
    if bar.get_height() >= 0:
        ax.text(
            bar.get_x() + bar.get_width()/2,
            bar.get_height() + 0.5,
            f"{bar.get_height():.2f}%",
            ha='center',
            fontsize=10
        )
    else:
        ax.text(
            bar.get_x() + bar.get_width()/2,
            bar.get_height() - 1.0,
            f"{bar.get_height():.2f}%",
            ha='center',
            fontsize=10
        )

plt.title('Percentage Improvement in Accuracy After Parameter Optimizatio
plt.xlabel('Model', fontsize=14)
plt.ylabel('Improvement (%)', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

Performing 10-fold cross-validation for Gaussian NB with best parameter  
S...

Average: AUC=0.7703 ± 0.0396  
Average: Accuracy=0.7377 ± 0.0363  
Average: F1 Score=0.6381 ± 0.0403

Performing 10-fold cross-validation for Multinomial NB with best parameter  
S...

Average: AUC=0.8660 ± 0.0281  
Average: Accuracy=0.8000 ± 0.0277  
Average: F1 Score=0.6273 ± 0.0495

Performing 10-fold cross-validation for Complement NB with best parameter  
S...

Average: AUC=0.8549 ± 0.0274  
Average: Accuracy=0.7851 ± 0.0343  
Average: F1 Score=0.6695 ± 0.0449

Performing 10-fold cross-validation for Bernoulli NB with best parameter  
S...

Average: AUC=0.8681 ± 0.0237  
Average: Accuracy=0.8096 ± 0.0229  
Average: F1 Score=0.7092 ± 0.0316

Performing 10-fold cross-validation for KNN with best parameters...

Average: AUC=0.7673 ± 0.0471  
Average: Accuracy=0.7263 ± 0.0260  
Average: F1 Score=0.5524 ± 0.0539

Performing 10-fold cross-validation for SVM with best parameters...

Average: AUC=0.8655 ± 0.0213  
Average: Accuracy=0.7956 ± 0.0390  
Average: F1 Score=0.6462 ± 0.0696

Performing 10-fold cross-validation for Logistic Regression with best parameters...

Average: AUC=0.8541 ± 0.0249  
Average: Accuracy=0.7825 ± 0.0277  
Average: F1 Score=0.6320 ± 0.0482



Performing 10-fold cross-validation for Random Forest with best parameter  
S...

Average: AUC=0.8466 ± 0.0334  
Average: Accuracy=0.7851 ± 0.0264  
Average: F1 Score=0.5831 ± 0.0627

/tmp/ipykernel\_368067/2330540877.py:127: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


ax = sns.barplot(



```
/tmp/ipykernel_368067/2330540877.py:224: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

```
ax = sns.barplot(
```



## Preprocessing

```
In [92]: import numpy as np

# Set style
sns.set(style="whitegrid")
plt.figure(figsize=(15, 10))

# Plot distribution of science_related tweets
plt.subplot(2, 2, 1)
sns.countplot(x='science_related', data=df)
plt.title('Distribution of Science Related Tweets')
plt.xlabel('Is Science Related')
plt.ylabel('Count')

# Plot distribution of science subtypes for science-related tweets
sci_df = df[df['science_related'] == 1]
plt.subplot(2, 2, 2)
subtypes = ['scientific_claim', 'scientific_reference', 'scientific_content']
sns.barplot(x=subtypes, y=[sci_df[col].sum() for col in subtypes])
plt.title('Distribution of Science Subtypes')
plt.xticks(rotation=45)
plt.tight_layout()

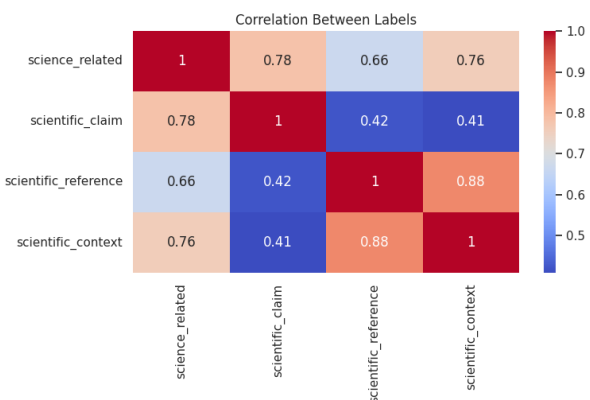

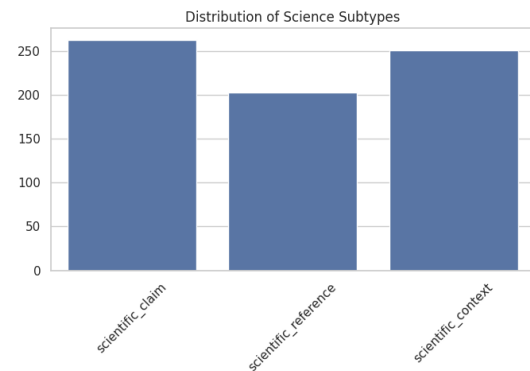

# Plot text length distribution
plt.subplot(2, 2, 3)
df['text_length'] = df['text'].apply(len)
sns.histplot(data=df, x='text_length', hue='science_related', bins=50, kde=True)
plt.title('Text Length Distribution by Class')
plt.xlabel('Text Length')
plt.ylabel('Frequency')

# Plot correlation between features
plt.subplot(2, 2, 4)
```



```
corr_cols = ['science_related', 'scientific_claim', 'scientific_reference']
sns.heatmap(df[corr_cols].corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Between Labels')

plt.tight_layout()
plt.show()
```



```
In [93]: import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
import emoji

# Download required NLTK resources
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('punkt')
nltk.download('punkt_tab')

def preprocess_text(text):
    # Convert to lowercase
    text = text.lower()

    # Demojize text
    text = emoji.demojize(text)

    # Remove URLs
    text = re.sub(r'(http\S+|www\S+)', '', text)

    # Remove mentions and hashtags execept for the word eurekamag
    text = re.sub(r'@\w+', '', text)
    text = re.sub(r'#eurekamag', 'eurekamag', text)
    text = re.sub(r'#\w+', '', text)
```

```

# Remove special characters and numbers
text = re.sub(r'^a-zA-Z\s', '', text)

# Tokenize
tokens = nltk.word_tokenize(text)

# Remove stopwords
stop_words = set(stopwords.words('english'))
tokens = [word for word in tokens if word not in stop_words]

# Lemmatize
lemmatizer = WordNetLemmatizer()
tokens = [lemmatizer.lemmatize(word) for word in tokens]

# Rejoin
return ' '.join(tokens)

# Apply preprocessing to the dataset
df['processed_text'] = df['text'].apply(preprocess_text)

# Compare before and after
comparison_df = pd.DataFrame({
    'Original': df['text'].head(5),
    'Preprocessed': df['processed_text'].head(5)
})
comparison_df

```

```

[nltk_data] Downloading package stopwords to /home/hurel/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /home/hurel/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package punkt to /home/hurel/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to /home/hurel/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!

```

Out[93]:

	Original	Preprocessed
0	Knees are a bit sore. i guess that's a sign th...	knee bit sore guess thats sign recent treadmil...
1	McDonald's breakfast stop then the gym 🏀👊	mcdonalds breakfast stop gym basketballflexedb...
2	Can any Gynecologist with Cancer Experience ex...	gynecologist cancer experience explain danger ...
3	Couch-lock highs lead to sleeping in the couch...	couchlock high lead sleeping couch got ta stop...
4	Does daily routine help prevent problems with ...	daily routine help prevent problem bipolar dis...

In [94]: # Visualize the impact of preprocessing

```

plt.figure(figsize=(12, 6))

# Text length before and after preprocessing
df['original_length'] = df['text'].apply(len)
df['processed_length'] = df['processed_text'].apply(len)

plt.subplot(1, 2, 1)

```

```
In [121]: # Compare different feature extraction methods
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_selection import SelectKBest, chi2, mutual_info_classif

# Create label for Task 1
y_task1 = df['science_related']

# 1. Count Vectorizer
count_vec = CountVectorizer(max_features=5000)
X_count = count_vec.fit_transform(df['processed_text'])

# 2. TF-IDF with more parameters
tfidf_vec = TfidfVectorizer(max_features=5000,
                             min_df=5,
                             max_df=0.8,
                             ngram_range=(1, 2))
X_tfidf = tfidf_vec.fit_transform(df['processed_text'])

# 3. TF-IDF with preprocessing already done
```

```

tfidf_processed = TfidfVectorizer(max_features=5000)
X_tfidf_processed = tfidf_processed.fit_transform(df['processed_text'])

# Compare feature extraction methods
print(f"Count Vectorizer Features: {X_count.shape}")
print(f"TF-IDF Vectorizer Features: {X_tfidf.shape}")
print(f"TF-IDF on Preprocessed Text Features: {X_tfidf_processed.shape}")

# Feature selection using Chi-squared
selector_chi2 = SelectKBest(chi2, k=100)
X_chi2 = selector_chi2.fit_transform(X_tfidf, y_task1)

# Feature selection using Mutual Information
selector_mi = SelectKBest(mutual_info_classif, k=100)
X_mi = selector_mi.fit_transform(X_tfidf, y_task1)

print(f"\nFeatures after Chi-squared selection: {X_chi2.shape}")
print(f"Features after Mutual Information selection: {X_mi.shape}")

# Get and visualize the most important features
chi2_selected_indices = selector_chi2.get_support(indices=True)
mi_selected_indices = selector_mi.get_support(indices=True)

chi2_feature_names = np.array(tfidf_vec.get_feature_names_out())[chi2_selected_indices]
mi_feature_names = np.array(tfidf_vec.get_feature_names_out())[mi_selected_indices]

# Plot top 20 features by importance
plt.figure(figsize=(16, 8))

plt.subplot(1, 2, 1)
chi2_scores = selector_chi2.scores_[chi2_selected_indices]
chi2_features_df = pd.DataFrame({'Feature': chi2_feature_names, 'Score': chi2_scores})
chi2_features_df = chi2_features_df.sort_values('Score', ascending=False)
sns.barplot(x='Score', y='Feature', data=chi2_features_df)
plt.title('Top 20 Features - Chi-squared')

plt.subplot(1, 2, 2)
mi_scores = selector_mi.scores_[mi_selected_indices]
mi_features_df = pd.DataFrame({'Feature': mi_feature_names, 'Score': mi_scores})
mi_features_df = mi_features_df.sort_values('Score', ascending=False)
sns.barplot(x='Score', y='Feature', data=mi_features_df)
plt.title('Top 20 Features - Mutual Information')

plt.tight_layout()
plt.show()

# We'll use the TF-IDF on preprocessed text for subsequent modeling
X_selected = X_tfidf

```

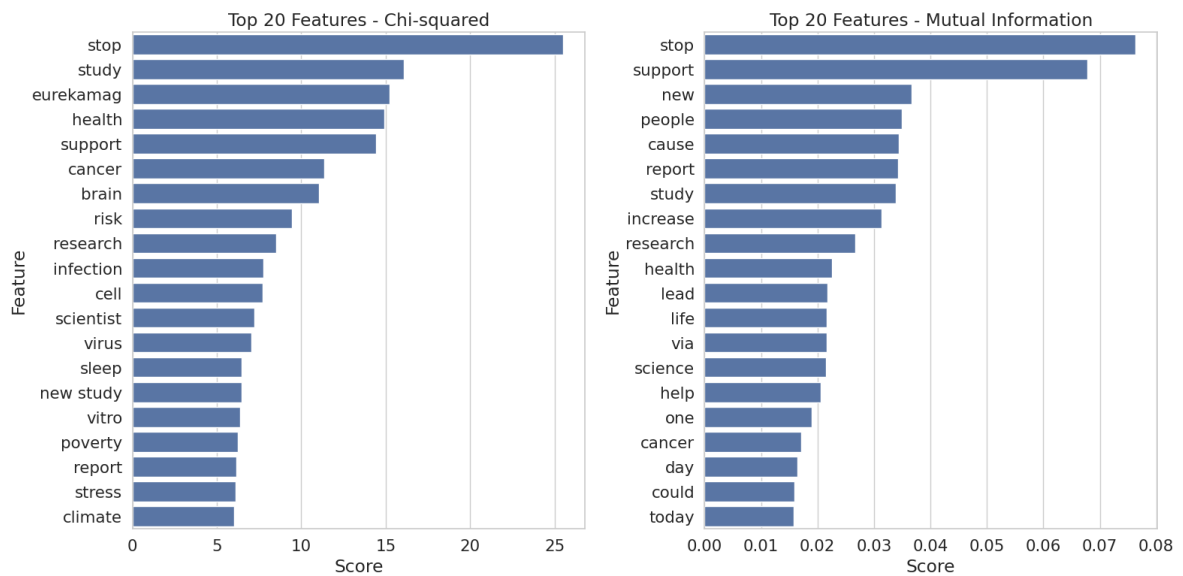
Count Vectorizer Features: (1140, 5000)

TF-IDF Vectorizer Features: (1140, 465)

TF-IDF on Preprocessed Text Features: (1140, 5000)

Features after Chi-squared selection: (1140, 100)

Features after Mutual Information selection: (1140, 100)



```
In [97]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB, MultinomialNB, ComplementNB,
from sklearn.metrics import classification_report, confusion_matrix, accu

# Function to evaluate model performance
def evaluate_model(model, X_train, X_test, y_train, y_test, model_name):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)
    report = classification_report(y_test, y_pred, output_dict=True)
    cm = confusion_matrix(y_test, y_pred)

    print(f"Model: {model_name}")
    print(f"Accuracy: {accuracy:.4f}")
    print(classification_report(y_test, y_pred))

    return {
        'model_name': model_name,
        'accuracy': accuracy,
        'report': report,
        'confusion_matrix': cm,
        'y_pred': y_pred
    }

# Get a dense version of our features for Gaussian NB
X_dense = X_selected.toarray()

# Split data for Task 1
X_train_task1, X_test_task1, y_train_task1, y_test_task1 = train_test_spl
    X_dense, df['task1_label'], test_size=0.2, random_state=42
)

# Define models with their best parameters based on grid search results
nb_models = {
    'Gaussian NB': GaussianNB(var_smoothing=0.0152), #ample value - updat
    'Multinomial NB': MultinomialNB(alpha=0.4641, fit_prior=True), # Exa
    'Complement NB': ComplementNB(alpha=2.8480, fit_prior=True), # Examp
    'Bernoulli NB': BernoulliNB(alpha=0.17475, fit_prior=True), # Examp
    'KNN': KNeighborsClassifier(n_neighbors=7, weights='uniform', metric=
    'SVM': SVC(C=1.4174742, kernel='rbf', gamma='scale', probability=True
    'Logistic Regression': LogisticRegression(C=np.float64(4.328761281083
```

```
    'Random Forest': RandomForestClassifier(n_estimators=200, max_depth=N
}

task1_results = {}
print("Task 1: Science Related Classification\n" + "="*40)
for name, model in nb_models.items():
    task1_results[name] = evaluate_model(
        model, X_train_task1, X_test_task1, y_train_task1, y_test_task1,
    )
    print("\n")
```

## Task 1: Science Related Classification

=====

Model: Gaussian NB

Accuracy: 0.6798

	precision	recall	f1-score	support
0	0.80	0.66	0.73	146
1	0.54	0.71	0.61	82
accuracy			0.68	228
macro avg	0.67	0.69	0.67	228
weighted avg	0.71	0.68	0.69	228

Model: Multinomial NB

Accuracy: 0.7412

	precision	recall	f1-score	support
0	0.75	0.89	0.82	146
1	0.71	0.48	0.57	82
accuracy			0.74	228
macro avg	0.73	0.68	0.69	228
weighted avg	0.74	0.74	0.73	228

Model: Complement NB

Accuracy: 0.7500

	precision	recall	f1-score	support
0	0.82	0.78	0.80	146
1	0.64	0.70	0.67	82
accuracy			0.75	228
macro avg	0.73	0.74	0.73	228
weighted avg	0.76	0.75	0.75	228

Model: Bernoulli NB

Accuracy: 0.7500

	precision	recall	f1-score	support
0	0.78	0.86	0.81	146
1	0.69	0.56	0.62	82
accuracy			0.75	228
macro avg	0.73	0.71	0.72	228
weighted avg	0.74	0.75	0.74	228

Model: KNN

Accuracy: 0.4737

	precision	recall	f1-score	support
0	0.91	0.20	0.33	146
1	0.40	0.96	0.57	82

accuracy			0.47	228
macro avg	0.65	0.58	0.45	228
weighted avg	0.73	0.47	0.41	228

Model: SVM

Accuracy: 0.7763

	precision	recall	f1-score	support
0	0.78	0.90	0.84	146
1	0.75	0.56	0.64	82

accuracy			0.78	228
macro avg	0.77	0.73	0.74	228
weighted avg	0.77	0.78	0.77	228

Model: Logistic Regression

Accuracy: 0.7675

	precision	recall	f1-score	support
0	0.79	0.88	0.83	146
1	0.72	0.57	0.64	82

accuracy			0.77	228
macro avg	0.75	0.72	0.73	228
weighted avg	0.76	0.77	0.76	228

Model: Random Forest

Accuracy: 0.7588

	precision	recall	f1-score	support
0	0.83	0.79	0.81	146
1	0.65	0.71	0.68	82

accuracy			0.76	228
macro avg	0.74	0.75	0.74	228
weighted avg	0.76	0.76	0.76	228

```
In [98]: # Using preprocessed text for feature extraction
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score

# Create TF-IDF features from preprocessed text
tfidf_vec = TfidfVectorizer(max_features=5000,
                           min_df=5,
                           max_df=0.8,
                           ngram_range=(1, 2))
X_tfidf_processed = tfidf_vec.fit_transform(df['processed_text'])

# Feature selection using Chi-squared
selector_chi2 = SelectKBest(chi2, k=100)
```



```
X_chi2_processed = selector_chi2.fit_transform(X_tfidf_processed, y_task1)

# Get a dense version of our features for models that require it
X_processed_dense = X_tfidf_processed.toarray()

# Initialize k-fold cross validation
k_folds = 10
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)

# Define models with their best parameters based on grid search results
best_models = {
    'Gaussian NB': GaussianNB(var_smoothing=0.0152),
    'Multinomial NB': MultinomialNB(alpha=0.4641, fit_prior=True),
    'Complement NB': ComplementNB(alpha=2.8480, fit_prior=True),
    'Bernoulli NB': BernoulliNB(alpha=0.17475, fit_prior=True),
    'KNN': KNeighborsClassifier(n_neighbors=7, weights='uniform', metric=
    'SVM': SVC(C=1.4174742, kernel='rbf', gamma='scale', probability=True
    'Logistic Regression': LogisticRegression(C=np.float64(4.328761281083
    'Random Forest': RandomForestClassifier(n_estimators=200, max_depth=N
}

# Dictionary to store results
best_cv_results = {}

# Perform k-fold cross validation for each model
for model_name, model in best_models.items():
    print(f"Performing {k_folds}-fold cross-validation for {model_name} w

    # Initialize lists to store performance metrics for each fold
    fold_accuracy = []
    fold_precision = []
    fold_recall = []
    fold_f1 = []
    fold_auc = []

    # For each fold
    for fold, (train_idx, test_idx) in enumerate(skf.split(X_processed_de
        # Split data
        X_train_fold, X_test_fold = X_processed_dense[train_idx], X_proce
        y_train_fold, y_test_fold = y_task1.iloc[train_idx], y_task1.iloc

        # Train model
        model.fit(X_train_fold, y_train_fold)

        # Make predictions
        y_pred_fold = model.predict(X_test_fold)

        # Calculate metrics
        acc = accuracy_score(y_test_fold, y_pred_fold)
        prec = precision_score(y_test_fold, y_pred_fold, zero_division=0)
        rec = recall_score(y_test_fold, y_pred_fold, zero_division=0)
        f1 = f1_score(y_test_fold, y_pred_fold, zero_division=0)

        # For AUC, we need probability estimates
        try:
            if hasattr(model, "predict_proba"):
                y_prob = model.predict_proba(X_test_fold)[: , 1]
                auc_score = roc_auc_score(y_test_fold, y_prob)
                fold_auc.append(auc_score)
        except:
```

```

        pass

        fold_accuracy.append(acc)
        fold_precision.append(prec)
        fold_recall.append(rec)
        fold_f1.append(f1)

    # Store average metrics and standard deviations
    best_cv_results[model_name] = {
        'accuracy': {
            'mean': np.mean(fold_accuracy),
            'std': np.std(fold_accuracy)
        },
        'precision': {
            'mean': np.mean(fold_precision),
            'std': np.std(fold_precision)
        },
        'recall': {
            'mean': np.mean(fold_recall),
            'std': np.std(fold_recall)
        },
        'f1': {
            'mean': np.mean(fold_f1),
            'std': np.std(fold_f1)
        }
    }

    if fold_auc:
        best_cv_results[model_name]['auc'] = {
            'mean': np.mean(fold_auc),
            'std': np.std(fold_auc)
        }
        print(f" Average: AUC={best_cv_results[model_name]['auc']['mean']}")

    print(f" Average: Accuracy={best_cv_results[model_name]['accuracy']['mean']}")
    print(f" Average: F1 Score={best_cv_results[model_name]['f1']['mean']}")
    print()

# Create DataFrame for visualization
best_results_df = pd.DataFrame({
    'Model': [],
    'Metric': [],
    'Mean': [],
    'Std': []
})

for model_name in best_cv_results:
    for metric in ['accuracy', 'precision', 'recall', 'f1']:
        best_results_df = pd.concat([best_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': [metric.capitalize()],
            'Mean': [best_cv_results[model_name][metric]['mean']],
            'Std': [best_cv_results[model_name][metric]['std']]
        })], ignore_index=True)
    if 'auc' in best_cv_results[model_name]:
        best_results_df = pd.concat([best_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': ['AUC'],
            'Mean': [best_cv_results[model_name]['auc']['mean']],
            'Std': [best_cv_results[model_name]['auc']['std']]
        })], ignore_index=True)

```

```

    }]], ignore_index=True)

# Sort models by accuracy
best_model_order = best_results_df[best_results_df['Metric'] == 'Accuracy']

# Create plots
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")

# Create bar plot for accuracy
ax = sns.barplot(
    data=best_results_df[best_results_df['Metric'] == 'Accuracy'],
    x='Model',
    y='Mean',
    order=best_model_order,
    palette='Blues_d'
)

# Add error bars
for i, model in enumerate(best_model_order):
    row = best_results_df[(best_results_df['Model'] == model) & (best_res
ax.errorbar(
    i, row['Mean'], yerr=row['Std'],
    fmt='o', color='black', elinewidth=2, capsize=6
)

# Add value labels on top of bars
for i, bar in enumerate(ax.patches):
    ax.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height() + 0.01,
        f"{bar.get_height():.3f}",
        ha='center',
        fontsize=10
    )

plt.title(f'Model Accuracy Comparison with Preprocessed Text ({k_folds}-F
plt.xlabel('Model', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.ylim([0.5, 1.0])
plt.tight_layout()
plt.show()

# Create a grouped bar chart for all metrics
plt.figure(figsize=(15, 10))
sns.set_style("whitegrid")

# Create grouped bar plot
ax = sns.barplot(
    data=best_results_df[best_results_df['Metric'] != 'AUC'],
    x='Model',
    y='Mean',
    hue='Metric',
    order=best_model_order,
    palette='Set2'
)

plt.title(f'Model Performance with Preprocessed Text ({k_folds}-Fold CV)'
plt.xlabel('Model', fontsize=14)

```

```
plt.ylabel('Score', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Metric', loc='lower right')
plt.ylim([0.5, 1.0])
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

Performing 10-fold cross-validation for Gaussian NB with best parameters on preprocessed text...

Average: AUC=0.7498 ± 0.0386  
Average: Accuracy=0.7211 ± 0.0332  
Average: F1 Score=0.6283 ± 0.0560

Performing 10-fold cross-validation for Multinomial NB with best parameters on preprocessed text...

Average: AUC=0.8442 ± 0.0334  
Average: Accuracy=0.7895 ± 0.0370  
Average: F1 Score=0.6251 ± 0.0685

Performing 10-fold cross-validation for Complement NB with best parameters on preprocessed text...

Average: AUC=0.8550 ± 0.0328  
Average: Accuracy=0.7675 ± 0.0310  
Average: F1 Score=0.6750 ± 0.0425

Performing 10-fold cross-validation for Bernoulli NB with best parameters on preprocessed text...

Average: AUC=0.8421 ± 0.0246  
Average: Accuracy=0.7851 ± 0.0270  
Average: F1 Score=0.6510 ± 0.0474

Performing 10-fold cross-validation for KNN with best parameters on preprocessed text...

Average: AUC=0.6058 ± 0.0533  
Average: Accuracy=0.5886 ± 0.0745  
Average: F1 Score=0.4461 ± 0.1151

Performing 10-fold cross-validation for SVM with best parameters on preprocessed text...

Average: AUC=0.8463 ± 0.0323  
Average: Accuracy=0.7982 ± 0.0229  
Average: F1 Score=0.6534 ± 0.0402

Performing 10-fold cross-validation for Logistic Regression with best parameters on preprocessed text...

Average: AUC=0.8347 ± 0.0318  
Average: Accuracy=0.7667 ± 0.0289  
Average: F1 Score=0.6182 ± 0.0508


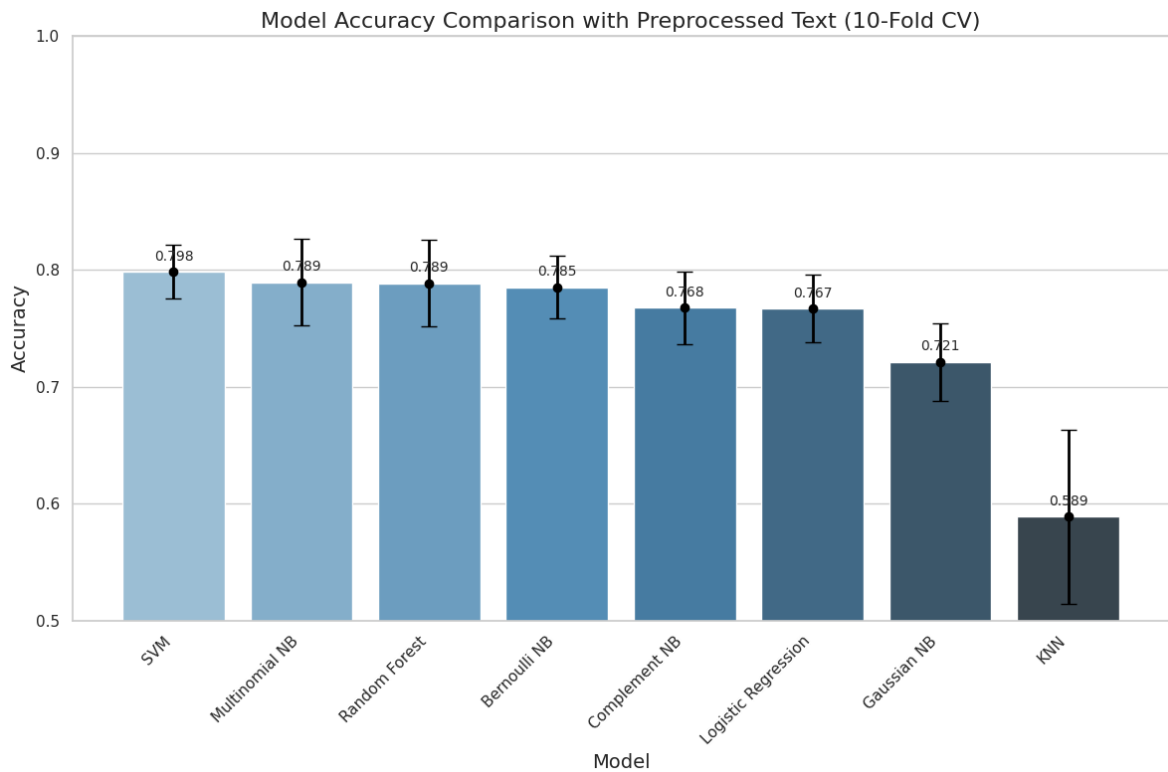
Performing 10-fold cross-validation for Random Forest with best parameters on preprocessed text...

Average: AUC=0.8573 ± 0.0307  
Average: Accuracy=0.7886 ± 0.0373  
Average: F1 Score=0.6906 ± 0.0609

```
/tmp/ipykernel_368067/2418365623.py:144: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be remove  
d in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for  
the same effect.
```

```
ax = sns.barplot(
```



```
In [99]: import seaborn as sns  
import pandas as pd  
import numpy as np  
from sklearn.metrics import confusion_matrix, roc_curve, auc
```

```

# Compare performance of SVM with preprocessed data and Bernoulli NB with
import matplotlib.pyplot as plt

# Create a DataFrame for comparison
comparison_data = [
    {
        'Model': 'Bernoulli NB (Raw Text)',
        'Accuracy': best_cv_results['Bernoulli NB']['accuracy']['mean'],
        'Precision': best_cv_results['Bernoulli NB']['precision']['mean'],
        'Recall': best_cv_results['Bernoulli NB']['recall']['mean'],
        'F1-Score': best_cv_results['Bernoulli NB']['f1']['mean']
    },
    {
        'Model': 'SVM (Preprocessed Text)',
        'Accuracy': best_cv_results['SVM']['accuracy']['mean'],
        'Precision': best_cv_results['SVM']['precision']['mean'],
        'Recall': best_cv_results['SVM']['recall']['mean'],
        'F1-Score': best_cv_results['SVM']['f1']['mean']
    }
]

# Create DataFrame
compare_df = pd.DataFrame(comparison_data)

# Reshape data for visualization
compare_melted = pd.melt(compare_df, id_vars='Model',
                          value_vars=['Accuracy', 'Precision', 'Recall', 'F1-Score'],
                          var_name='Metric', value_name='Score')

# Plot comparison
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")

# Create grouped bar chart
ax = sns.barplot(data=compare_melted, x='Metric', y='Score', hue='Model',
                 order=['Accuracy', 'Precision', 'Recall', 'F1-Score'])

plt.title('Performance Comparison: Bernoulli NB (Raw) vs. SVM (Preprocessed)')
plt.xlabel('Metric', fontsize=14)
plt.ylabel('Score', fontsize=14)
plt.ylim([0.4, 1.0])
plt.grid(True, linestyle='--', alpha=0.7)
plt.legend(title='Model')

# Add value labels on bars
for p in ax.patches:
    ax.annotate(f'{p.get_height():.3f}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha = 'center', va = 'bottom',
                fontsize=10)

plt.tight_layout()
plt.show()

# Statistical significance test (if applicable)
print("\nPerformance Summary:")
print(compare_df.set_index('Model'))

# Analyze differences
diff_accuracy = abs(comparison_data[0]['Accuracy'] - comparison_data[1]['Accuracy'])
diff_f1 = abs(comparison_data[0]['F1-Score'] - comparison_data[1]['F1-Score'])

```


```

print(f"\nDifference in Accuracy: {diff_accuracy:.3f}")
print(f"Difference in F1-Score: {diff_f1:.3f}")

# Print conclusions
if comparison_data[0]['Accuracy'] > comparison_data[1]['Accuracy']:
    winner = "Bernoulli NB with raw text"
else:
    winner = "SVM with preprocessed text"

print(f"\nConclusion: {winner} performs better in terms of accuracy.")
print("This suggests that " +
      ("text preprocessing is beneficial for this classification task."
       if winner == "SVM with preprocessed text"
       else "raw text features might contain important signals that are l

```



Performance Summary:

	Accuracy	Precision	Recall	F1-Score
Model				
Bernoulli NB (Raw Text)	0.785088	0.704412	0.613158	0.651009
SVM (Preprocessed Text)	0.798246	0.758331	0.581081	0.653448

Difference in Accuracy: 0.013

Difference in F1-Score: 0.002

Conclusion: SVM with preprocessed text performs better in terms of accuracy.  
 This suggests that text preprocessing is beneficial for this classification task.

```
In [100]... best_model = best_models['SVM']
```

```
In [101]... from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
import numpy as np

import matplotlib.pyplot as plt
```

```
# Get predictions from the best model on the test set
y_pred = best_model.predict(X_test_task1)

# Calculate the confusion matrix
cm = confusion_matrix(y_test_task1, y_pred)


# Create a prettier visualization of the confusion matrix
plt.figure(figsize=(10, 8))
sns.set(font_scale=1.4)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                  xticklabels=['Not Science', 'Science'],
                  yticklabels=['Not Science', 'Science'])

# Add labels, title and ticks
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.title('Confusion Matrix - SVM Model')

# Calculate metrics
report = classification_report(y_test_task1, y_pred, output_dict=True)
accuracy = (cm[0, 0] + cm[1, 1]) / np.sum(cm)
precision = report['1']['precision']
recall = report['1']['recall']
f1 = report['1']['f1-score']

# Display metrics on the plot
plt.figtext(0.5, 0.01, f'Accuracy: {accuracy:.4f} | Precision: {precision:.4f} | Recall: {recall:.4f} | F1-Score: {f1:.4f}',
            ha='center', fontsize=12, bbox={'facecolor': 'lightblue', 'alpha': 0.5})

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```



Accuracy: 0.9605 | Precision: 0.9398 | Recall: 0.9512 | F1-Score: 0.9455



In [102... `%pip install imbalanced-learn`

```
# Import required modules
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler, SMOTE
from collections import Counter
```

Requirement already satisfied: imbalanced-learn in ./venv/lib/python3.10/site-packages (0.13.0)  
 Requirement already satisfied: numpy<3,>=1.24.3 in ./venv/lib/python3.10/site-packages (from imbalanced-learn) (2.2.5)  
 Requirement already satisfied: scipy<2,>=1.10.1 in ./venv/lib/python3.10/site-packages (from imbalanced-learn) (1.15.2)  
 Requirement already satisfied: scikit-learn<2,>=1.3.2 in ./venv/lib/python3.10/site-packages (from imbalanced-learn) (1.6.1)  
 Requirement already satisfied: sklearn-compat<1,>=0.1 in ./venv/lib/python3.10/site-packages (from imbalanced-learn) (0.1.3)  
 Requirement already satisfied: joblib<2,>=1.1.1 in ./venv/lib/python3.10/site-packages (from imbalanced-learn) (1.4.2)  
 Requirement already satisfied: threadpoolctl<4,>=2.0.0 in ./venv/lib/python3.10/site-packages (from imbalanced-learn) (3.6.0)  
 Note: you may need to restart the kernel to use updated packages.

In [103... `# For Task 1 (science_related classification)`

```
def apply_undersampling(X, y):
    undersampler = RandomUnderSampler(random_state=42)
    X_resampled, y_resampled = undersampler.fit_resample(X, y)
    print(f"Original distribution: {Counter(y)}")
    print(f"Distribution after undersampling: {Counter(y_resampled)}")
    return X_resampled, y_resampled
```

`# Apply to Task 1`

```
X_train_task1_under, y_train_task1_under = apply_undersampling(X_train_ta
```

Original distribution: Counter({0: 619, 1: 293})

Distribution after undersampling: Counter({0: 293, 1: 293})

In [104... `# Visualize effect of sampling`

```
def plot_sampling_comparison(original_y, resampled_y, title):
    plt.figure(figsize=(12, 5))
```


```
    plt.subplot(1, 2, 1)
    sns.countplot(x=original_y)
    plt.title('Original Distribution')
```

```
    plt.subplot(1, 2, 2)
    sns.countplot(x=resampled_y)
    plt.title('After Resampling')
```

```
    plt.suptitle(title)
    plt.tight_layout()
    plt.show()
```

`# Example for Task 1`

```
plot_sampling_comparison(y_train_task1, y_train_task1_under, 'Effect of U
```



```
In [105... # Retrain the best model (SVM) with undersampled data
best_model = best_models['SVM']

# Evaluate on the original test data for fair comparison
best_model.fit(X_train_task1_under, y_train_task1_under)
y_pred_under = best_model.predict(X_test_task1)
accuracy_under = accuracy_score(y_test_task1, y_pred_under)
report_under = classification_report(y_test_task1, y_pred_under, output_d
cm_under = confusion_matrix(y_test_task1, y_pred_under)

# Perform k-fold cross validation with undersampled data
# Initialize StratifiedKFold
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)

# Store evaluation metrics
fold_accuracy_under = []
fold_precision_under = []
fold_recall_under = []
fold_f1_under = []
fold_auc_under = []

# Get a balanced dataset
X_balanced = X_train_task1_under
y_balanced = y_train_task1_under

# Perform cross-validation
for fold, (train_idx, test_idx) in enumerate(skf.split(X_balanced, y_bala
X_train_fold, X_test_fold = X_balanced[train_idx], X_balanced[test_id
y_train_fold, y_test_fold = y_balanced.iloc[train_idx], y_balanced.il

best_model.fit(X_train_fold, y_train_fold)
y_pred_fold = best_model.predict(X_test_fold)

# Calculate metrics
acc = accuracy_score(y_test_fold, y_pred_fold)
prec = precision_score(y_test_fold, y_pred_fold, zero_division=0)
rec = recall_score(y_test_fold, y_pred_fold, zero_division=0)
f1 = f1_score(y_test_fold, y_pred_fold, zero_division=0)

# For AUC
if hasattr(best_model, "predict_proba"):
    y_prob = best_model.predict_proba(X_test_fold)[: , 1]
    auc_score = roc_auc_score(y_test_fold, y_prob)
    fold_auc_under.append(auc_score)
```

```

fold_accuracy_under.append(acc)
fold_precision_under.append(prec)
fold_recall_under.append(rec)
fold_f1_under.append(f1)

# Compare performance metrics before and after undersampling
comparison_data = {
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score', 'AUC'],
    'Original': [
        best_cv_results['SVM']['accuracy']['mean'],
        best_cv_results['SVM']['precision']['mean'],
        best_cv_results['SVM']['recall']['mean'],
        best_cv_results['SVM']['f1']['mean'],
        best_cv_results['SVM']['auc']['mean'] if 'auc' in best_cv_results
    ],
    'Undersampled': [
        np.mean(fold_accuracy_under),
        np.mean(fold_precision_under),
        np.mean(fold_recall_under),
        np.mean(fold_f1_under),
        np.mean(fold_auc_under) if fold_auc_under else None
    ]
}

comparison_df = pd.DataFrame(comparison_data)

# Create visualizations
plt.figure(figsize=(16, 10))

# Bar chart comparison
plt.subplot(2, 2, 1)
comparison_melted = pd.melt(comparison_df, id_vars=['Metric'],
                             value_vars=['Original', 'Undersampled'],
                             var_name='Data Balance', value_name='Score')

# Filter out any None values
comparison_melted = comparison_melted.dropna()

sns.barplot(x='Metric', y='Score', hue='Data Balance', data=comparison_me
plt.title('SVM Performance: Original vs. Undersampled', fontsize=14)
plt.ylim(0.5, 1.0)
plt.grid(True, linestyle='--', alpha=0.7)

# Confusion matrix for undersampled model
plt.subplot(2, 2, 2)
sns.heatmap(cm_under, annot=True, fmt='d', cmap='Blues',
             xticklabels=['Not Science', 'Science'],
             yticklabels=['Not Science', 'Science'])
plt.title('Confusion Matrix (Undersampled)', fontsize=14)
plt.ylabel('True Label')
plt.xlabel('Predicted Label')

# ROC Curve comparison if available
if hasattr(best_model, "predict_proba"):
    plt.subplot(2, 2, 3)
    # For original model
    y_prob_orig = best_model.fit(X_train_task1, y_train_task1).predict_pr
    fpr_orig, tpr_orig, _ = roc_curve(y_test_task1, y_prob_orig)
    auc_orig = auc(fpr_orig, tpr_orig)

```

```
# For undersampled model
best_model.fit(X_train_task1_under, y_train_task1_under)
y_prob_under = best_model.predict_proba(X_test_task1)[:, 1]
fpr_under, tpr_under, _ = roc_curve(y_test_task1, y_prob_under)
auc_under = auc(fpr_under, tpr_under)

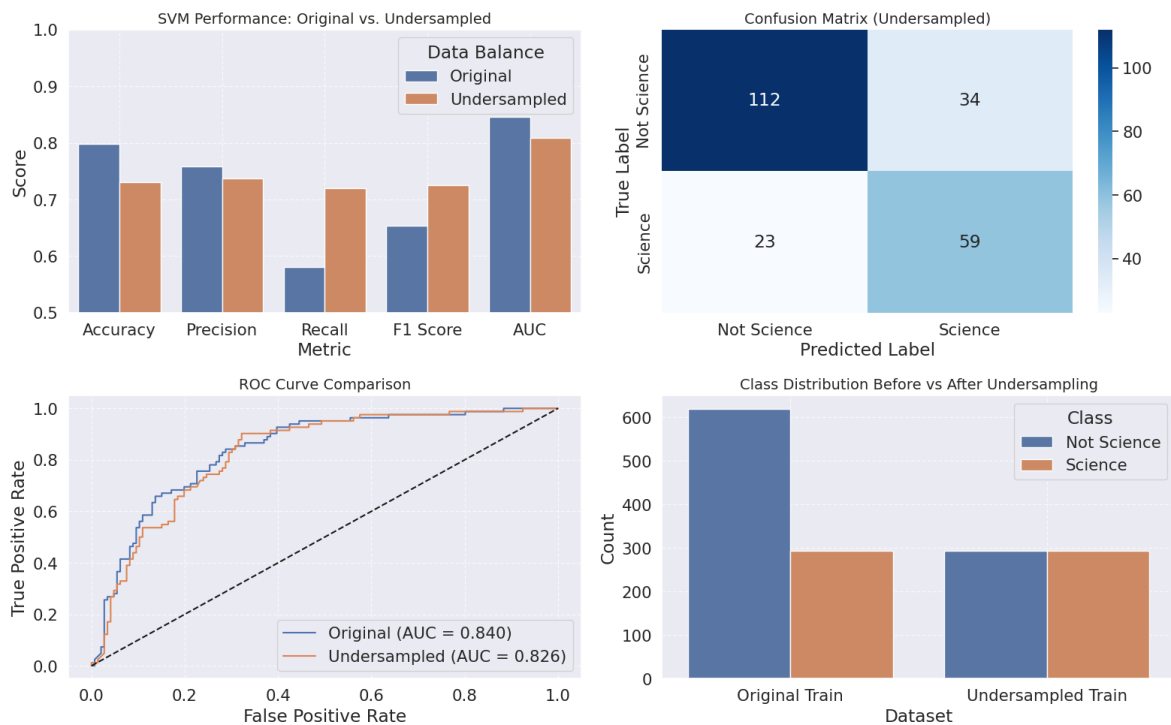
# Plot ROC curves
plt.plot(fpr_orig, tpr_orig, label=f'Original (AUC = {auc_orig:.3f})')
plt.plot(fpr_under, tpr_under, label=f'Undersampled (AUC = {auc_under:.3f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.title('ROC Curve Comparison', fontsize=14)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)

# Class distribution before and after undersampling
plt.subplot(2, 2, 4)
class_dist = pd.DataFrame({
    'Dataset': ['Original Train'] * 2 + ['Undersampled Train'] * 2,
    'Class': ['Not Science', 'Science'] * 2,
    'Count': [
        sum(y_train_task1 == 0),
        sum(y_train_task1 == 1),
        sum(y_train_task1_under == 0),
        sum(y_train_task1_under == 1)
    ]
})
sns.barplot(x='Dataset', y='Count', hue='Class', data=class_dist)
plt.title('Class Distribution Before vs After Undersampling', fontsize=14)
plt.grid(True, linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()

# Print summary
print("\nSVM Performance Comparison Summary:")
print(comparison_df.set_index('Metric'))
print("\nOriginal Class Distribution:", Counter(y_train_task1))
print("Undersampled Class Distribution:", Counter(y_train_task1_under))

# Calculate improvement percentages
improvement_df = pd.DataFrame({
    'Metric': comparison_df['Metric'],
    'Improvement (%)': [
        ((comparison_df['Undersampled'][i] - comparison_df['Original'][i]) /
         comparison_df['Original'][i] * 100) if comparison_df['Original'][i] is not None
        else None
        for i in range(len(comparison_df))
    ]
})
print("\nRelative Improvement After Undersampling:")
print(improvement_df.set_index('Metric'))
```



### SVM Performance Comparison Summary:

	Original	Undersampled
Metric		
Accuracy	0.798246	0.730479
Precision	0.758331	0.737308
Recall	0.581081	0.719885
F1 Score	0.653448	0.725460
AUC	0.846289	0.809384

Original Class Distribution: Counter({0: 619, 1: 293})

Undersampled Class Distribution: Counter({0: 293, 1: 293})

### Relative Improvement After Undersampling:

Metric	Improvement (%)
Accuracy	-8.489412
Precision	-2.772259
Recall	23.887196
F1 Score	11.020415
AUC	-4.360885

```
In [106]: # TASK 2: Scientific Claim/Reference Classification

# Since Task 2 is only applicable to science-related tweets, we need to u
print("Task 2: Scientific Claim/Reference Classification\n" + "="*50)

# Get a dense version of features for science-related tweets
# Convert the pandas Series to a numpy array with .values
X_dense_sci = X_selected[(df['science_related'] == 1).values].toarray()
y_task2 = df_sci['task2_label']

# Split data for Task 2
X_train_task2, X_test_task2, y_train_task2, y_test_task2 = train_test_spl
    X_dense_sci, y_task2, test_size=0.2, random_state=42
)

# Define the best models from Task 1 (you can modify these based on Task
task2_models = {
```

```
'Gaussian NB': GaussianNB(var_smoothing=0.0152),
'Multinomial NB': MultinomialNB(alpha=0.4641, fit_prior=True),
'Complement NB': ComplementNB(alpha=2.8480, fit_prior=True),
'Bernoulli NB': BernoulliNB(alpha=0.17475, fit_prior=True),
'KNN': KNeighborsClassifier(n_neighbors=7, weights='uniform', metric=
'SVM': SVC(C=1.4174742, kernel='rbf', gamma='scale', probability=True)
'Logistic Regression': LogisticRegression(C=np.float64(4.328761281083
'Random Forest': RandomForestClassifier(n_estimators=200, max_depth=N
}

# Train and evaluate models for Task 2
task2_results = {}
for name, model in task2_models.items():
    task2_results[name] = evaluate_model(
        model, X_train_task2, X_test_task2, y_train_task2, y_test_task2,
    )
print("\n")
```

## Task 2: Scientific Claim/Reference Classification

=====

Model: Gaussian NB

Accuracy: 0.8533

	precision	recall	f1-score	support
0	0.30	0.43	0.35	7
1	0.94	0.90	0.92	68
accuracy			0.85	75
macro avg	0.62	0.66	0.64	75
weighted avg	0.88	0.85	0.86	75

Model: Multinomial NB

Accuracy: 0.9067

	precision	recall	f1-score	support
0	0.00	0.00	0.00	7
1	0.91	1.00	0.95	68
accuracy			0.91	75
macro avg	0.45	0.50	0.48	75
weighted avg	0.82	0.91	0.86	75

Model: Complement NB

Accuracy: 0.7200

	precision	recall	f1-score	support
0	0.18	0.57	0.28	7
1	0.94	0.74	0.83	68
accuracy			0.72	75
macro avg	0.56	0.65	0.55	75
weighted avg	0.87	0.72	0.78	75

Model: Bernoulli NB

Accuracy: 0.8533

	precision	recall	f1-score	support
0	0.00	0.00	0.00	7
1	0.90	0.94	0.92	68
accuracy			0.85	75
macro avg	0.45	0.47	0.46	75
weighted avg	0.82	0.85	0.83	75

[illegible]



Model: KNN

Accuracy: 0.9067

	precision	recall	f1-score	support
0	0.00	0.00	0.00	7
1	0.91	1.00	0.95	68
accuracy			0.91	75
macro avg	0.45	0.50	0.48	75
weighted avg	0.82	0.91	0.86	75

Model: SVM

Accuracy: 0.9200

	precision	recall	f1-score	support
0	1.00	0.14	0.25	7
1	0.92	1.00	0.96	68
accuracy			0.92	75
macro avg	0.96	0.57	0.60	75
weighted avg	0.93	0.92	0.89	75

Model: Logistic Regression

Accuracy: 0.9200

	precision	recall	f1-score	support
0	1.00	0.14	0.25	7
1	0.92	1.00	0.96	68
accuracy			0.92	75
macro avg	0.96	0.57	0.60	75
weighted avg	0.93	0.92	0.89	75

Model: Random Forest

Accuracy: 0.9200

	precision	recall	f1-score	support
0	1.00	0.14	0.25	7
1	0.92	1.00	0.96	68
accuracy			0.92	75
macro avg	0.96	0.57	0.60	75
weighted avg	0.93	0.92	0.89	75

```
In [107... # Cross-validation for Task 2
print("Task 2: Cross-Validation\n" + "="*40)

# Initialize k-fold cross validation
k_folds = 10
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)
```

```
# Dictionary to store CV results
task2_cv_results = {}

# Perform k-fold cross validation for each model
for model_name, model in task2_models.items():
    print(f"Performing {k_folds}-fold cross-validation for {model_name}..

    # Initialize lists to store performance metrics for each fold
    fold_accuracy = []
    fold_precision = []
    fold_recall = []
    fold_f1 = []
    fold_auc = []

    # For each fold
    for fold, (train_idx, test_idx) in enumerate(skf.split(X_dense_sci, y
        # Split data
        X_train_fold, X_test_fold = X_dense_sci[train_idx], X_dense_sci[t
        y_train_fold, y_test_fold = y_task2.iloc[train_idx], y_task2.iloc

        # Train model
        model.fit(X_train_fold, y_train_fold)

        # Make predictions
        y_pred_fold = model.predict(X_test_fold)

        # Calculate metrics
        acc = accuracy_score(y_test_fold, y_pred_fold)
        prec = precision_score(y_test_fold, y_pred_fold, zero_division=0)
        rec = recall_score(y_test_fold, y_pred_fold, zero_division=0)
        f1 = f1_score(y_test_fold, y_pred_fold, zero_division=0)

        # For AUC, we need probability estimates
        try:
            if hasattr(model, "predict_proba"):
                y_prob = model.predict_proba(X_test_fold)[: , 1]
                auc_score = roc_auc_score(y_test_fold, y_prob)
                fold_auc.append(auc_score)
        except:
            pass

        fold_accuracy.append(acc)
        fold_precision.append(prec)
        fold_recall.append(rec)
        fold_f1.append(f1)

    # Store average metrics and standard deviations
    task2_cv_results[model_name] = {
        'accuracy': {
            'mean': np.mean(fold_accuracy),
            'std': np.std(fold_accuracy)
        },
        'precision': {
            'mean': np.mean(fold_precision),
            'std': np.std(fold_precision)
        },
        'recall': {
            'mean': np.mean(fold_recall),
            'std': np.std(fold_recall)
        },
    },
```

```

        'f1': {
            'mean': np.mean(fold_f1),
            'std': np.std(fold_f1)
        }
    }

    if fold_auc:
        task2_cv_results[model_name]['auc'] = {
            'mean': np.mean(fold_auc),
            'std': np.std(fold_auc)
        }
        print(f"   Average: AUC={task2_cv_results[model_name]['auc']['mean']

print(f"   Average: Accuracy={task2_cv_results[model_name]['accuracy']}
print(f"   Average: F1 Score={task2_cv_results[model_name]['f1']['mean']
print()

```

## Task 2: Cross-Validation

=====

Performing 10-fold cross-validation for Gaussian NB...

Average: AUC=0.5660 ± 0.1187  
 Average: Accuracy=0.8217 ± 0.0635  
 Average: F1 Score=0.8988 ± 0.0375

Performing 10-fold cross-validation for Multinomial NB...

Average: AUC=0.7806 ± 0.0932  
 Average: Accuracy=0.9148 ± 0.0152  
 Average: F1 Score=0.9554 ± 0.0081

Performing 10-fold cross-validation for Complement NB...

Average: AUC=0.7468 ± 0.0979  
 Average: Accuracy=0.7151 ± 0.0768  
 Average: F1 Score=0.8205 ± 0.0532

Performing 10-fold cross-validation for Bernoulli NB...

Average: AUC=0.7860 ± 0.0991  
 Average: Accuracy=0.8935 ± 0.0332  
 Average: F1 Score=0.9423 ± 0.0184

Performing 10-fold cross-validation for KNN...

Average: AUC=0.5324 ± 0.0660  
 Average: Accuracy=0.9094 ± 0.0126  
 Average: F1 Score=0.9525 ± 0.0069

Performing 10-fold cross-validation for SVM...

Average: AUC=0.7151 ± 0.1219  
 Average: Accuracy=0.9175 ± 0.0179  
 Average: F1 Score=0.9567 ± 0.0095

Performing 10-fold cross-validation for Logistic Regression...

Average: AUC=0.7019 ± 0.1686  
 Average: Accuracy=0.9068 ± 0.0242  
 Average: F1 Score=0.9503 ± 0.0131

Performing 10-fold cross-validation for Random Forest...

Average: AUC=0.7275 ± 0.1402  
 Average: Accuracy=0.9148 ± 0.0152  
 Average: F1 Score=0.9554 ± 0.0081

```

In [108... # Visualize Task 2 cross-validation results
# Create DataFrame for visualization
task2_results_df = pd.DataFrame({
    'Model': [],
    'Metric': [],
    'Mean': [],
    'Std': []
})

for model_name in task2_cv_results:
    for metric in ['accuracy', 'precision', 'recall', 'f1']:
        task2_results_df = pd.concat([task2_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': [metric.capitalize()],
            'Mean': [task2_cv_results[model_name][metric]['mean']],
            'Std': [task2_cv_results[model_name][metric]['std']]
        })], ignore_index=True)
    if 'auc' in task2_cv_results[model_name]:
        task2_results_df = pd.concat([task2_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': ['AUC'],
            'Mean': [task2_cv_results[model_name]['auc']['mean']],
            'Std': [task2_cv_results[model_name]['auc']['std']]
        })], ignore_index=True)

# Sort models by accuracy
task2_model_order = task2_results_df[task2_results_df['Metric'] == 'Accuracy']

# Create accuracy bar plot
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")

# Create bar plot for accuracy
ax = sns.barplot(
    data=task2_results_df[task2_results_df['Metric'] == 'Accuracy'],
    x='Model',
    y='Mean',
    order=task2_model_order,
    palette='Blues_d'
)

# Add error bars
for i, model in enumerate(task2_model_order):
    row = task2_results_df[(task2_results_df['Model'] == model) & (task2_
ax.errorbar(
    i, row['Mean'], yerr=row['Std'],
    fmt='o', color='black', elinewidth=2, capsize=6
)

# Add value labels on top of bars
for i, bar in enumerate(ax.patches):
    ax.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height() + 0.01,
        f"{bar.get_height():.3f}",
        ha='center',
        fontsize=10
    )

plt.title(f'Model Accuracy Comparison for Task 2 with {k_folds}-Fold Cross-Validation')

```

```
plt.xlabel('Model', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.ylim([0.5, 1.0])
plt.tight_layout()
plt.show()

# Create a grouped bar chart for all metrics
plt.figure(figsize=(15, 10))
sns.set_style("whitegrid")


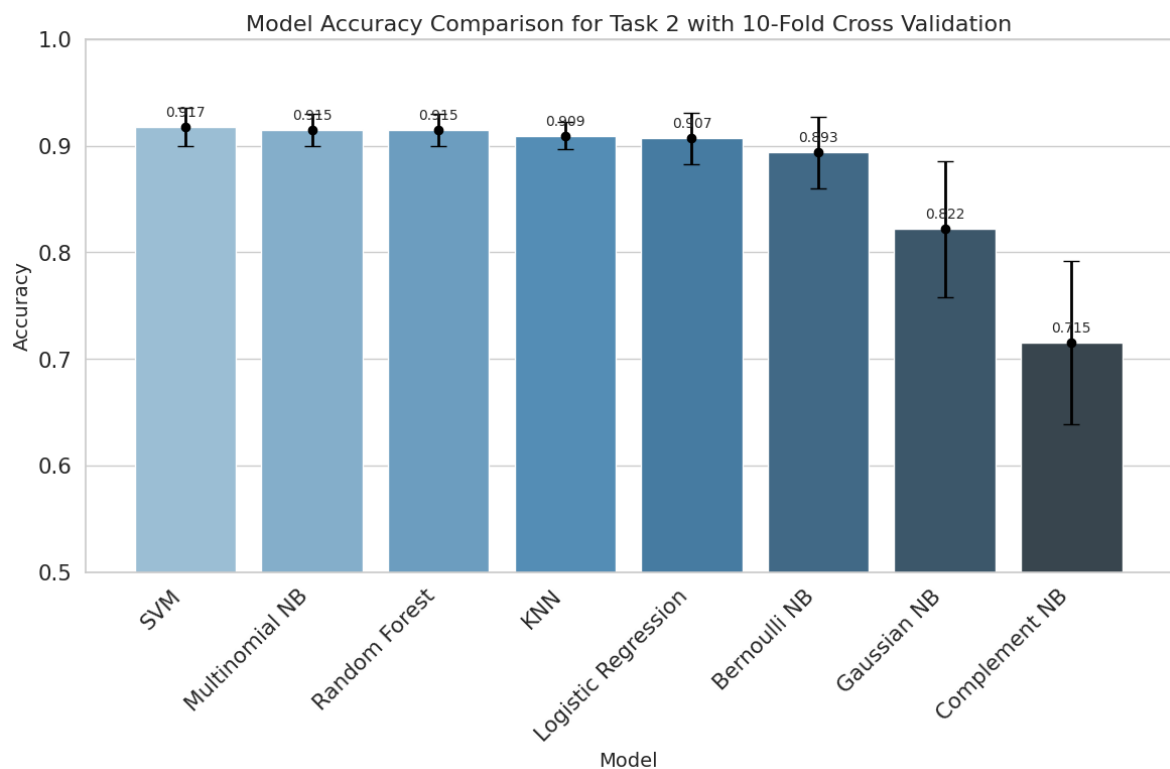
# Create grouped bar plot
ax = sns.barplot(
    data=task2_results_df[task2_results_df['Metric'] != 'AUC'],
    x='Model',
    y='Mean',
    hue='Metric',
    order=task2_model_order,
    palette='Set2'
)

plt.title(f'Task 2: Model Performance Metrics Comparison ({k_folds}-Fold
plt.xlabel('Model', fontsize=14)
plt.ylabel('Score', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Metric', loc='lower right')
plt.ylim([0.5, 1.0])
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

/tmp/ipykernel\_368067/1744879752.py:34: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(
```



```
In [109... # Hyperparameter optimization for the best model on Task 2
# Let's assume SVM was the best model (update based on your results)
best_task2_model_name = 'SVM' # Replace with the actual best model from
best_task2_model = task2_models[best_task2_model_name]

print(f"Performing Grid Search for {best_task2_model_name} on Task 2...")

# Define the parameter grid for the best model
if best_task2_model_name == 'SVM':
    param_grid = {
        'C': np.logspace(-3, 3, 10),
        'gamma': ['scale', 'auto'] + list(np.logspace(-3, 3, 5)),
        'kernel': ['rbf', 'linear']
```

```

    }
elif best_task2_model_name == 'Random Forest':
    param_grid = {
        'n_estimators': [50, 100, 200],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    }
elif best_task2_model_name.endswith('NB'):
    if best_task2_model_name == 'Gaussian NB':
        param_grid = {
            'var_smoothing': np.logspace(-10, 0, 11)
        }
    else:
        param_grid = {
            'alpha': np.logspace(-3, 3, 10),
            'fit_prior': [True, False]
        }
else:
    # Default grid for other models
    param_grid = {
        'C': np.logspace(-3, 3, 10)
    }

# Create and fit GridSearchCV
grid_search_task2 = GridSearchCV(
    estimator=best_task2_model,
    param_grid=param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

grid_search_task2.fit(X_train_task2, y_train_task2)

# Display the best parameters and score
print(f"Best Parameters: {grid_search_task2.best_params_}")
print(f"Best Cross-Validation Score: {grid_search_task2.best_score_:.4f}")

# Evaluate on test set
best_model_task2 = grid_search_task2.best_estimator_
y_pred_task2 = best_model_task2.predict(X_test_task2)
accuracy_task2 = accuracy_score(y_test_task2, y_pred_task2)
print(f"Test Accuracy with Best Parameters: {accuracy_task2:.4f}")
print("\nClassification Report:")
print(classification_report(y_test_task2, y_pred_task2))

# Plot confusion matrix
plt.figure(figsize=(10, 8))
cm = confusion_matrix(y_test_task2, y_pred_task2)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Claim/Ref', 'Claim/Ref'],
            yticklabels=['No Claim/Ref', 'Claim/Ref'])
plt.title(f'Confusion Matrix - Task 2 ({best_task2_model_name})')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.tight_layout()
plt.show()

```

Performing Grid Search for SVM on Task 2...

Fitting 5 folds for each of 140 candidates, totalling 700 fits


Best Parameters: {'C': np.float64(215.44346900318823), 'gamma': 'auto', 'kernel': 'rbf'}

Best Cross-Validation Score: 0.9167

Test Accuracy with Best Parameters: 0.9200

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.14	0.25	7
1	0.92	1.00	0.96	68
accuracy			0.92	75
macro avg	0.96	0.57	0.60	75
weighted avg	0.93	0.92	0.89	75



```
In [110... # Check for class imbalance in Task 2
print("Task 2 class distribution:")
print(y_task2.value_counts())
print(y_task2.value_counts(normalize=True).round(3) * 100, '%')

# Apply class balancing for Task 2 if needed
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from sklearn.base import clone # Add this import

# Let's try SMOTE for handling imbalance
```



```
smote = SMOTE(random_state=42)
X_train_task2_balanced, y_train_task2_balanced = smote.fit_resample(X_train_task2, y_train_task2)

print("\nClass distribution after SMOTE balancing:")
print(pd.Series(y_train_task2_balanced).value_counts())

# Train the best model with balanced data
best_model_balanced = clone(best_model_task2)
best_model_balanced.fit(X_train_task2_balanced, y_train_task2_balanced)

# Evaluate on the test set
y_pred_balanced = best_model_balanced.predict(X_test_task2)
accuracy_balanced = accuracy_score(y_test_task2, y_pred_balanced)
print(f"\nTest Accuracy with Balanced Training Data: {accuracy_balanced:.4f}")
print("\nClassification Report with Balanced Training Data:")
print(classification_report(y_test_task2, y_pred_balanced))

# Compare balanced vs unbalanced training
plt.figure(figsize=(12, 6))
metrics = ['Accuracy', 'Precision', 'Recall', 'F1-Score']
unbalanced_scores = [
    accuracy_task2,
    precision_score(y_test_task2, y_pred_task2),
    recall_score(y_test_task2, y_pred_task2),
    f1_score(y_test_task2, y_pred_task2)
]
balanced_scores = [
    accuracy_balanced,
    precision_score(y_test_task2, y_pred_balanced),
    recall_score(y_test_task2, y_pred_balanced),
    f1_score(y_test_task2, y_pred_balanced)
]

x = np.arange(len(metrics))
width = 0.35

plt.bar(x - width/2, unbalanced_scores, width, label='Unbalanced Training')
plt.bar(x + width/2, balanced_scores, width, label='Balanced Training')
plt.xticks(x, metrics)
plt.ylabel('Score')
plt.title('Impact of Class Balancing on Task 2 Performance')
plt.legend()
plt.ylim(0, 1)
plt.tight_layout()
plt.show()
```

Task 2 class distribution:

task2\_label

1 342

0 33

Name: count, dtype: int64

task2\_label

1 91.2

0 8.8

Name: proportion, dtype: float64 %

Class distribution after SMOTE balancing:

task2\_label

1 274


0 274

Name: count, dtype: int64

Test Accuracy with Balanced Training Data: 0.8800

Classification Report with Balanced Training Data:

	precision	recall	f1-score	support
0	0.25	0.14	0.18	7
1	0.92	0.96	0.94	68
accuracy			0.88	75
macro avg	0.58	0.55	0.56	75
weighted avg	0.85	0.88	0.86	75



```
In [111]: # TASK 3: Scientific Type Classification
print("Task 3: Scientific Type Classification\n" + "="*50)

# We already have the preprocessed features for science-related tweets (X)
# And we have the labels in df_sci['task3_label'] (with values 0, 1, 2 for
y_task3 = df_sci['task3_label']

# Split data for Task 3
X_train_task3, X_test_task3, y_train_task3, y_test_task3 = train_test_split(
    X_dense_sci, y_task3, test_size=0.2, random_state=42
)

# Use the same models as in Task 2
task3_models = {
```

```
'Gaussian NB': GaussianNB(var_smoothing=0.0152),
'Multinomial NB': MultinomialNB(alpha=0.4641, fit_prior=True),
'Complement NB': ComplementNB(alpha=2.8480, fit_prior=True),
'Bernoulli NB': BernoulliNB(alpha=0.17475, fit_prior=True),
'KNN': KNeighborsClassifier(n_neighbors=7, weights='uniform', metric=
'SVM': SVC(C=1.4174742, kernel='rbf', gamma='scale', probability=True)
'Logistic Regression': LogisticRegression(C=np.float64(4.328761281083
'Random Forest': RandomForestClassifier(n_estimators=200, max_depth=N
}

# Train and evaluate models for Task 3
task3_results = {}
for name, model in task3_models.items():
    task3_results[name] = evaluate_model(
        model, X_train_task3, X_test_task3, y_train_task3, y_test_task3,
    )
print("\n")
```

## Task 3: Scientific Type Classification

Model: Gaussian NB

Accuracy: 0.7600

	precision	recall	f1-score	support
0	0.86	0.89	0.87	54
1	0.67	0.43	0.52	14
2	0.30	0.43	0.35	7
accuracy			0.76	75
macro avg	0.61	0.58	0.58	75
weighted avg	0.77	0.76	0.76	75

Model: Multinomial NB

Accuracy: 0.7600

	precision	recall	f1-score	support
0	0.76	1.00	0.86	54
1	0.75	0.21	0.33	14
2	0.00	0.00	0.00	7
accuracy			0.76	75
macro avg	0.50	0.40	0.40	75
weighted avg	0.69	0.76	0.68	75

Model: Complement NB

Accuracy: 0.6000

	precision	recall	f1-score	support
0	0.78	0.72	0.75	54
1	0.31	0.29	0.30	14
2	0.17	0.29	0.21	7
accuracy			0.60	75
macro avg	0.42	0.43	0.42	75
weighted avg	0.63	0.60	0.61	75

Model: Bernoulli NB

Accuracy: 0.7067

	precision	recall	f1-score	support
0	0.78	0.91	0.84	54
1	0.43	0.21	0.29	14
2	0.20	0.14	0.17	7
accuracy			0.71	75
macro avg	0.47	0.42	0.43	75
weighted avg	0.66	0.71	0.67	75

Model: KNN

Accuracy: 0.7200

	precision	recall	f1-score	support
0	0.73	0.96	0.83	54
1	0.50	0.14	0.22	14
2	0.00	0.00	0.00	7
accuracy			0.72	75
macro avg	0.41	0.37	0.35	75
weighted avg	0.62	0.72	0.64	75

Model: SVM

Accuracy: 0.7733

	precision	recall	f1-score	support
0	0.79	1.00	0.89	54
1	0.60	0.21	0.32	14
2	0.50	0.14	0.22	7
accuracy			0.77	75
macro avg	0.63	0.45	0.47	75
weighted avg	0.73	0.77	0.72	75

```

/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/hurel/Documents/repo/projet-ml/.venv/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Model: Logistic Regression

Accuracy: 0.7067

	precision	recall	f1-score	support
0	0.77	0.89	0.83	54
1	0.33	0.29	0.31	14
2	1.00	0.14	0.25	7
accuracy			0.71	75
macro avg	0.70	0.44	0.46	75
weighted avg	0.71	0.71	0.68	75

Model: Random Forest

Accuracy: 0.6667

	precision	recall	f1-score	support
0	0.82	0.78	0.80	54
1	0.29	0.43	0.34	14
2	0.67	0.29	0.40	7
accuracy			0.67	75
macro avg	0.59	0.50	0.51	75
weighted avg	0.71	0.67	0.68	75

```
In [112]: # Cross-validation for Task 3
print("Task 3: Cross-Validation\n" + "="*40)

# Initialize k-fold cross validation
k_folds = 10
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)

# Dictionary to store CV results
task3_cv_results = {}

# Perform k-fold cross validation for each model
for model_name, model in task3_models.items():
    print(f"Performing {k_folds}-fold cross-validation for {model_name}..

    # Initialize lists to store performance metrics for each fold
    fold_accuracy = []
    fold_precision = []
    fold_recall = []
    fold_f1 = []

    # For each fold
    for fold, (train_idx, test_idx) in enumerate(skf.split(X_dense_sci, y
        # Split data
        X_train_fold, X_test_fold = X_dense_sci[train_idx], X_dense_sci[t
        y_train_fold, y_test_fold = y_task3.iloc[train_idx], y_task3.iloc

        # Train model
        model.fit(X_train_fold, y_train_fold)

        # Make predictions
        y_pred_fold = model.predict(X_test_fold)
```

```
# Calculate metrics
acc = accuracy_score(y_test_fold, y_pred_fold)
# For multi-class classification, use macro averaging
prec = precision_score(y_test_fold, y_pred_fold, average='macro', zero_div
rec = recall_score(y_test_fold, y_pred_fold, average='macro', zero_div
f1 = f1_score(y_test_fold, y_pred_fold, average='macro', zero_div

fold_accuracy.append(acc)
fold_precision.append(prec)
fold_recall.append(rec)
fold_f1.append(f1)

# Store average metrics and standard deviations
task3_cv_results[model_name] = {
    'accuracy': {
        'mean': np.mean(fold_accuracy),
        'std': np.std(fold_accuracy)
    },
    'precision': {
        'mean': np.mean(fold_precision),
        'std': np.std(fold_precision)
    },
    'recall': {
        'mean': np.mean(fold_recall),
        'std': np.std(fold_recall)
    },
    'f1': {
        'mean': np.mean(fold_f1),
        'std': np.std(fold_f1)
    }
}

print(f" Average: Accuracy={task3_cv_results[model_name]['accuracy']}
print(f" Average: F1 Score={task3_cv_results[model_name]['f1']['mean
print()
```



## Task 3: Cross-Validation

```

=====
Performing 10-fold cross-validation for Gaussian NB...
Average: Accuracy=0.6612 ± 0.0701
Average: F1 Score=0.4518 ± 0.0643

Performing 10-fold cross-validation for Multinomial NB...
Average: Accuracy=0.7280 ± 0.0599
Average: F1 Score=0.4004 ± 0.0922

Performing 10-fold cross-validation for Complement NB...
Average: Accuracy=0.6428 ± 0.0595
Average: F1 Score=0.4886 ± 0.0616

Performing 10-fold cross-validation for Bernoulli NB...
Average: Accuracy=0.6428 ± 0.0595
Average: F1 Score=0.4886 ± 0.0616

Performing 10-fold cross-validation for Bernoulli NB...
Average: Accuracy=0.7042 ± 0.0498
Average: F1 Score=0.4997 ± 0.0956

Performing 10-fold cross-validation for KNN...
Average: Accuracy=0.2983 ± 0.0671
Average: F1 Score=0.1982 ± 0.0542

Performing 10-fold cross-validation for SVM...
Average: Accuracy=0.7467 ± 0.0379
Average: F1 Score=0.4304 ± 0.0634

Performing 10-fold cross-validation for Logistic Regression...
Average: Accuracy=0.6962 ± 0.0395
Average: F1 Score=0.4445 ± 0.0834

Performing 10-fold cross-validation for Random Forest...
Average: Accuracy=0.6985 ± 0.0678
Average: F1 Score=0.4920 ± 0.1049

```

```

In [113]: # Visualize Task 3 cross-validation results
# Create DataFrame for visualization
task3_results_df = pd.DataFrame({
    'Model': [],
    'Metric': [],
    'Mean': [],
    'Std': []
})

for model_name in task3_cv_results:
    for metric in ['accuracy', 'precision', 'recall', 'f1']:
        task3_results_df = pd.concat([task3_results_df, pd.DataFrame({
            'Model': [model_name],
            'Metric': [metric.capitalize()],
            'Mean': [task3_cv_results[model_name][metric]['mean']],
            'Std': [task3_cv_results[model_name][metric]['std']]
        })], ignore_index=True)

# Sort models by accuracy
task3_model_order = task3_results_df[task3_results_df['Metric'] == 'Accur

```

```
# Create accuracy bar plot
plt.figure(figsize=(12, 8))
sns.set_style("whitegrid")

# Create bar plot for accuracy
ax = sns.barplot(
    data=task3_results_df[task3_results_df['Metric'] == 'Accuracy'],
    x='Model',
    y='Mean',
    order=task3_model_order,
    palette='Blues_d'
)

# Add error bars
for i, model in enumerate(task3_model_order):
    row = task3_results_df[(task3_results_df['Model'] == model) & (task3_
ax.errorbar(
    i, row['Mean'], yerr=row['Std'],
    fmt='o', color='black', elinewidth=2, capsize=6
)

# Add value labels on top of bars
for i, bar in enumerate(ax.patches):
    ax.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height() + 0.01,
        f"{bar.get_height():.3f}",
        ha='center',
        fontsize=10
    )

plt.title(f'Model Accuracy Comparison for Task 3 with {k_folds}-Fold Cross Validation')
plt.xlabel('Model', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.ylim([0.5, 1.0])
plt.tight_layout()
plt.show()

# Create a grouped bar chart for all metrics
plt.figure(figsize=(15, 10))
sns.set_style("whitegrid")


# Create grouped bar plot
ax = sns.barplot(
    data=task3_results_df,
    x='Model',
    y='Mean',
    hue='Metric',
    order=task3_model_order,
    palette='Set2'
)

plt.title(f'Task 3: Model Performance Metrics Comparison ({k_folds}-Fold Cross Validation)')
plt.xlabel('Model', fontsize=14)
plt.ylabel('Score', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Metric', loc='lower right')
plt.ylim([0.5, 1.0])
plt.grid(True, linestyle='--', alpha=0.7)
```

```
plt.tight_layout()  
plt.show()
```

/tmp/ipykernel\_368067/1263538886.py:27: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.



```
In [114]: # Choose the best model from cross-validation for Task 3  
best_task3_model_name = task3_model_order[0] # The model with highest accuracy  
best_task3_model = task3_models[best_task3_model_name]
```

```
# Hyperparameter optimization for the best model on Task 3
print(f"Performing Grid Search for {best_task3_model_name} on Task 3...")

# Define the parameter grid for the best model
if best_task3_model_name == 'SVM':
    param_grid = {
        'C': np.logspace(-3, 3, 7),
        'gamma': ['scale', 'auto'] + list(np.logspace(-3, 3, 5)),
        'kernel': ['rbf', 'linear']
    }
elif best_task3_model_name == 'Random Forest':
    param_grid = {
        'n_estimators': [50, 100, 200],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    }
elif best_task3_model_name.endswith('NB'):
    if best_task3_model_name == 'Gaussian NB':
        param_grid = {
            'var_smoothing': np.logspace(-10, 0, 11)
        }
    else:
        param_grid = {
            'alpha': np.logspace(-3, 3, 7),
            'fit_prior': [True, False]
        }
else:
    # Default grid for other models
    param_grid = {
        'C': np.logspace(-3, 3, 7)
    }

# Create and fit GridSearchCV
grid_search_task3 = GridSearchCV(
    estimator=best_task3_model,
    param_grid=param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

grid_search_task3.fit(X_train_task3, y_train_task3)

# Display the best parameters and score
print(f"Best Parameters: {grid_search_task3.best_params_}")
print(f"Best Cross-Validation Score: {grid_search_task3.best_score_:.4f}")

# Evaluate on test set
best_model_task3 = grid_search_task3.best_estimator_
y_pred_task3 = best_model_task3.predict(X_test_task3)
accuracy_task3 = accuracy_score(y_test_task3, y_pred_task3)
print(f"Test Accuracy with Best Parameters: {accuracy_task3:.4f}")
print("\nClassification Report:")
print(classification_report(y_test_task3, y_pred_task3))

# Plot confusion matrix
plt.figure(figsize=(10, 8))
cm = confusion_matrix(y_test_task3, y_pred_task3)
```

```

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Claim', 'Reference', 'Context'],
            yticklabels=['Claim', 'Reference', 'Context'])
plt.title(f'Confusion Matrix - Task 3 ({best_task3_model_name})')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.tight_layout()
plt.show()

```

Performing Grid Search for SVM on Task 3...

Fitting 5 folds for each of 98 candidates, totalling 490 fits


Best Parameters: {'C': np.float64(10.0), 'gamma': 'scale', 'kernel': 'rbf'}

Best Cross-Validation Score: 0.7433

Test Accuracy with Best Parameters: 0.7867

Classification Report:

	precision	recall	f1-score	support
0	0.82	1.00	0.90	54
1	0.57	0.29	0.38	14
2	0.50	0.14	0.22	7
accuracy			0.79	75
macro avg	0.63	0.48	0.50	75
weighted avg	0.74	0.79	0.74	75



```

In [115... # Check for class imbalance in Task 3
print("Task 3 class distribution:")
print(y_task3.value_counts())

```

```

print(y_task3.value_counts(normalize=True).round(3) * 100, '%')

# Apply SMOTE for handling class imbalance in Task 3
smote = SMOTE(random_state=42)
X_train_task3_balanced, y_train_task3_balanced = smote.fit_resample(X_tra

print("\nClass distribution after SMOTE balancing:")
print(pd.Series(y_train_task3_balanced).value_counts())

# Train the best model with balanced data
best_model_balanced = clone(best_model_task3)
best_model_balanced.fit(X_train_task3_balanced, y_train_task3_balanced)

# Evaluate on the test set
y_pred_balanced = best_model_balanced.predict(X_test_task3)
accuracy_balanced = accuracy_score(y_test_task3, y_pred_balanced)
print(f"\nTest Accuracy with Balanced Training Data: {accuracy_balanced:.")
print("\nClassification Report with Balanced Training Data:")
print(classification_report(y_test_task3, y_pred_balanced))

# Compare class-wise performance before and after balancing
original_report = classification_report(y_test_task3, y_pred_task3, output
balanced_report = classification_report(y_test_task3, y_pred_balanced, ou

# Create a DataFrame to compare per-class metrics
class_comparison = []
for class_label in ['0', '1', '2']: # Claim, Reference, Context
    for metric in ['precision', 'recall', 'f1-score']:
        class_comparison.append({
            'Class': ['Claim', 'Reference', 'Context'][int(class_label)],
            'Metric': metric.capitalize(),
            'Original': original_report[class_label][metric],
            'Balanced': balanced_report[class_label][metric],
            'Difference': balanced_report[class_label][metric] - original

        })

class_comp_df = pd.DataFrame(class_comparison)

# Plot class-wise performance comparison
plt.figure(figsize=(15, 10))

for i, cls in enumerate(['Claim', 'Reference', 'Context']):
    plt.subplot(1, 3, i+1)
    df_class = class_comp_df[class_comp_df['Class'] == cls]

    x = np.arange(len(df_class['Metric'].unique()))
    width = 0.35

    orig_scores = df_class['Original'].values
    bal_scores = df_class['Balanced'].values

    plt.bar(x - width/2, orig_scores, width, label='Original')
    plt.bar(x + width/2, bal_scores, width, label='Balanced')

    plt.title(f'{cls} Performance')
    plt.xticks(x, df_class['Metric'].unique())
    plt.ylim(0, 1)
    plt.legend()
    plt.grid(True, linestyle='--', alpha=0.7)

```

```
plt.suptitle('Impact of Class Balancing on Task 3 Performance by Class',  
plt.tight_layout(rect=[0, 0, 1, 0.95])  
plt.show()
```

Task 3 class distribution:

task3\_label

0 263

1 79

2 33

Name: count, dtype: int64

task3\_label

0 70.1

1 21.1

2 8.8

Name: proportion, dtype: float64 %

Class distribution after SMOTE balancing:

task3\_label

0 209

1 209

2 209


Name: count, dtype: int64

Test Accuracy with Balanced Training Data: 0.7733

Classification Report with Balanced Training Data:

	precision	recall	f1-score	support
0	0.79	1.00	0.89	54
1	0.60	0.21	0.32	14
2	0.50	0.14	0.22	7
accuracy			0.77	75
macro avg	0.63	0.45	0.47	75
weighted avg	0.73	0.77	0.72	75

Impact of Class Balancing on Task 3 Performance by Class



```

In [116... # Visualize confusion matrices for balanced vs unbalanced
plt.figure(figsize=(16, 7))

# Original model confusion matrix
plt.subplot(1, 2, 1)
cm_orig = confusion_matrix(y_test_task3, y_pred_task3)
sns.heatmap(cm_orig, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Claim', 'Reference', 'Context'],
            yticklabels=['Claim', 'Reference', 'Context'])
plt.title('Confusion Matrix - Original Model')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')

# Balanced model confusion matrix
plt.subplot(1, 2, 2)
cm_balanced = confusion_matrix(y_test_task3, y_pred_balanced)
sns.heatmap(cm_balanced, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Claim', 'Reference', 'Context'],
            yticklabels=['Claim', 'Reference', 'Context'])
plt.title('Confusion Matrix - Balanced Model')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')

plt.tight_layout()
plt.show()

```

